

# ➤ **Big Data**

## **Session 3: Column Stores and Coordination**

## Last week

- Distributed Infrastructures
- Google File System
- HDFS
- Apache Spark RDDs

## **At the end of this lecture, you will be able to:**

- Explain the benefits of column stores
- Describe approaches to schedule jobs in distributed infrastructures
- Outline how to manage distributed infrastructures

- Part 1:
  - Column stores (Introduction)
  - Big Table
  - HBase and Hive
- Part 2:
  - Job Scheduling
  - Coordination

# Relational Database

Attribute

Alan	17/03/1991	Koblenz	male	married
Steve	21/05/1999	Neuwied	male	single
Laura	13/06/1995	Andernach	female	single
Mei	30/04/1998	Koblenz	female	married

Tuple {

Relation

# Structured Query Language: SQL

- Declarative query language
- Multiple aspects of the language
  - Data definition language
    - Statements to create, modify tables and views
  - Data manipulation language
    - Statements to issue queries, insert, delete data

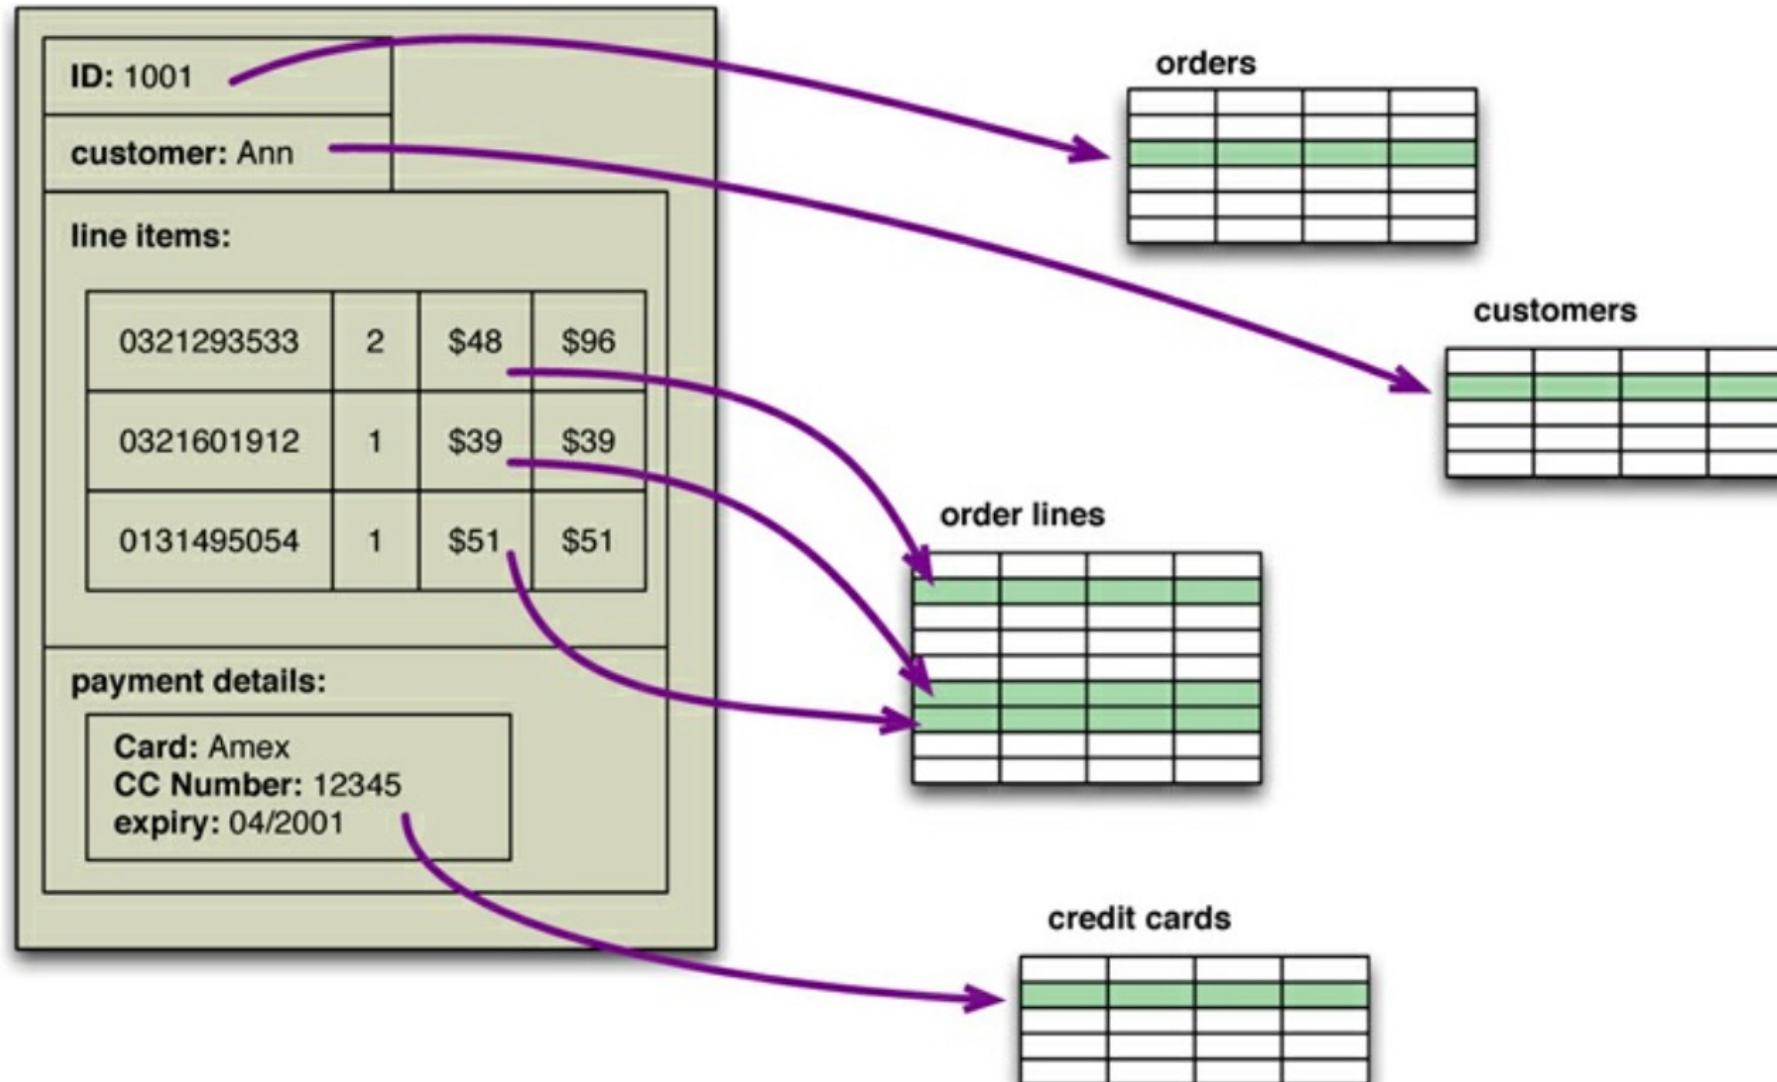
```
SELECT <attributes>  
FROM   <one or more relations>  
WHERE  <conditions>
```

# Example

## Building an e-commerce website

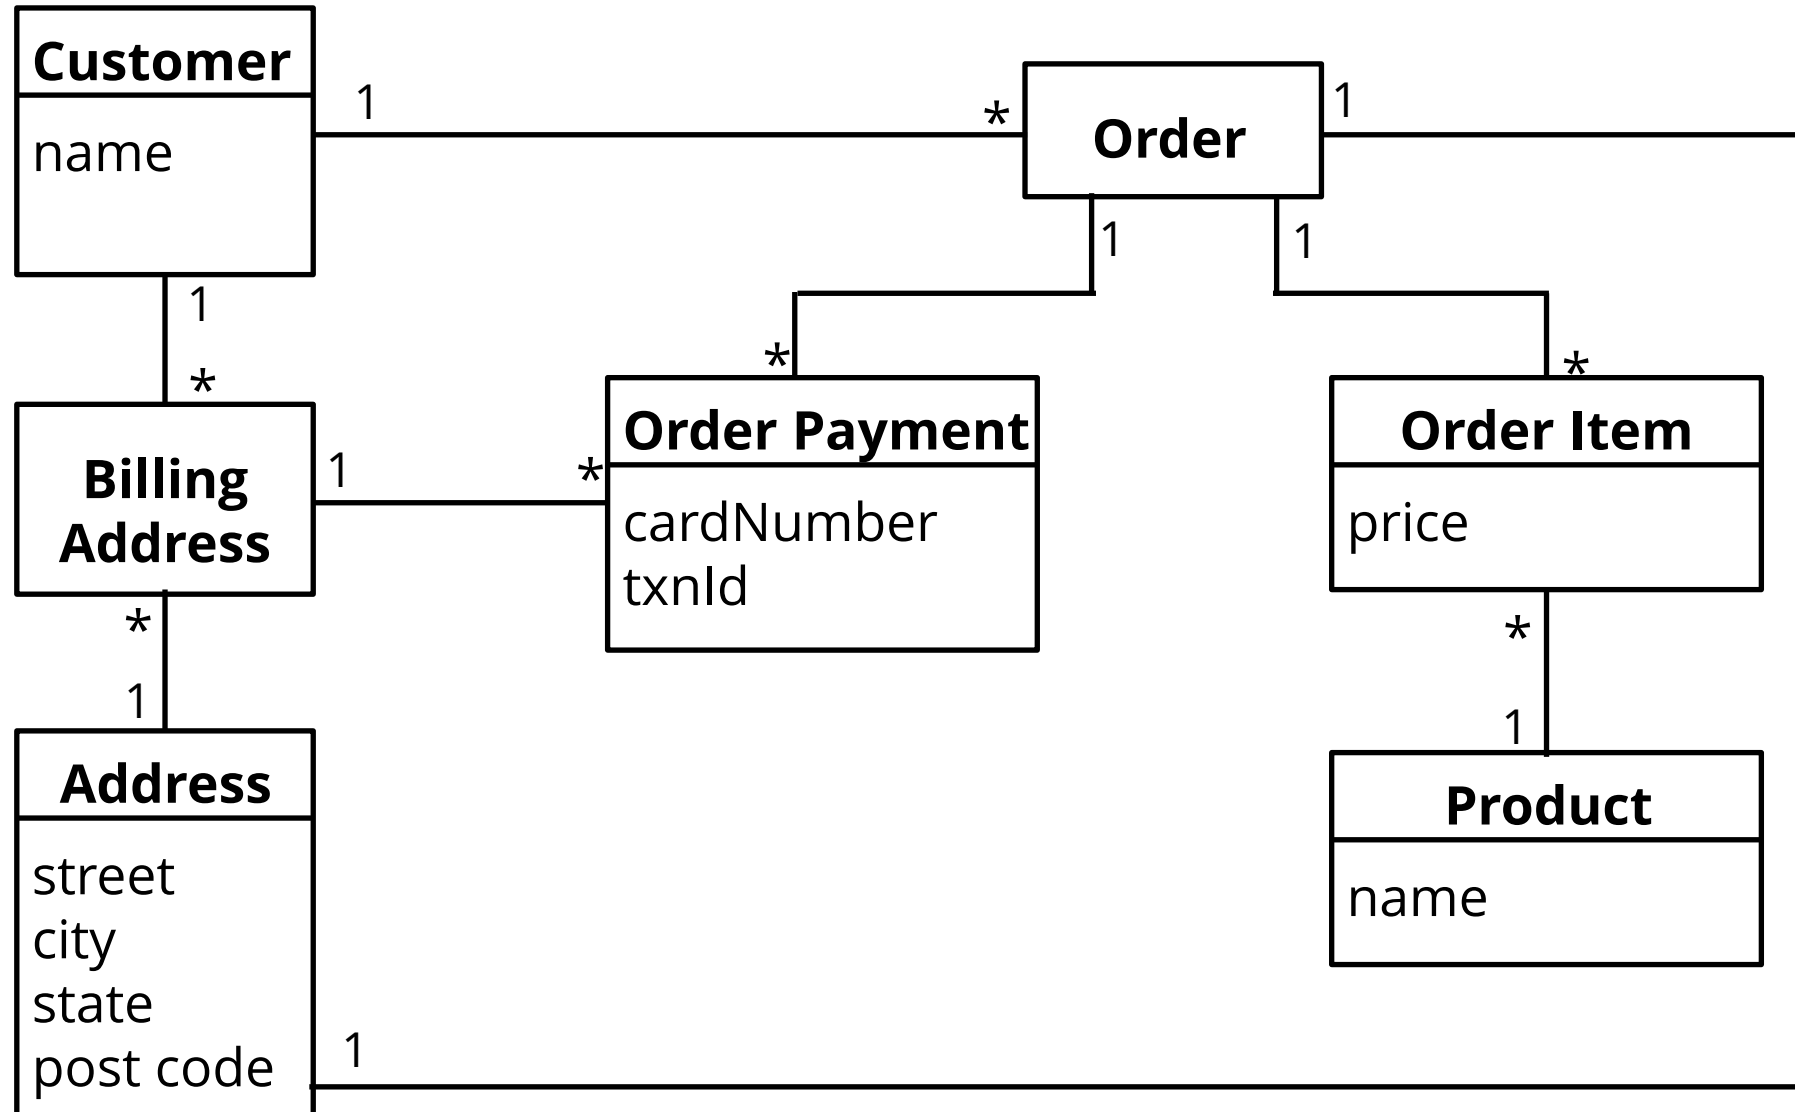
- Sell items directly to customers over the web
- Store information:
  - About users
  - Product catalog
  - Orders
  - Shipping addresses
  - Billing addresses
  - Payment data

# Relational model





# Relational Data Model



# Typical data using RDBMS data model

**Customer**

<b>Id</b>	<b>Name</b>
1	Frank

**Order**

<b>Id</b>	<b>CustomerId</b>	<b>ShippingAddressId</b>
99	1	77

**Product**

<b>Id</b>	<b>Name</b>
27	NoSQL Distilled

**BillingAddress**

<b>Id</b>	<b>CustomerId</b>	<b>AddressId</b>
55	1	77

**OrderItem**

<b>Id</b>	<b>OrderId</b>	<b>ProductId</b>	<b>Price</b>
100	99	27	32.45

**Address**

<b>Id</b>	<b>City</b>
77	Chicago

**OrderPayment**

<b>Id</b>	<b>OrderId</b>	<b>CardNumber</b>	<b>BillingAddressId</b>	<b>txnId</b>
33	99	1000-1000	55	abelif879rft

## Example: How to query these?

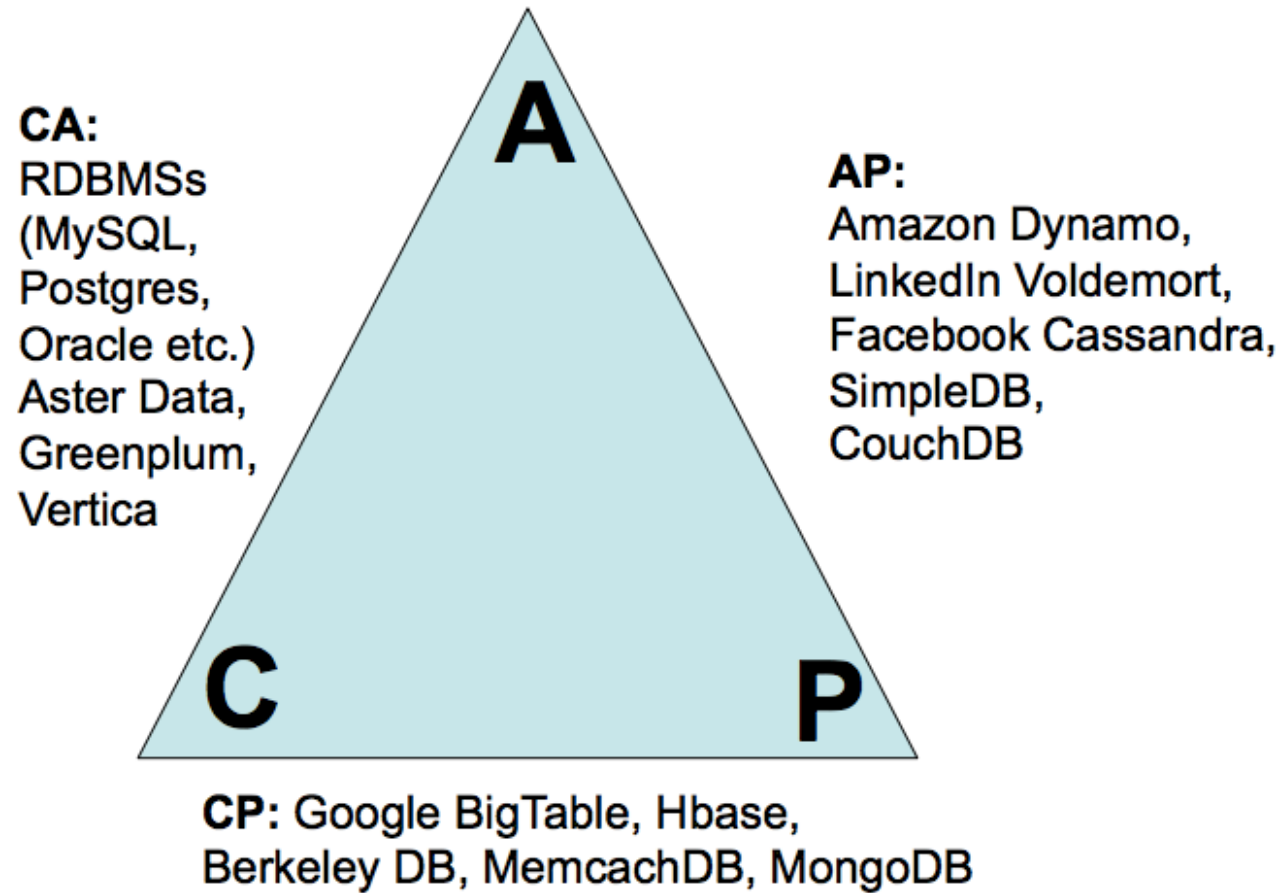
- What is the price of all orders in 2022?
- How many orders were payed by credit card in 2022?
- In how many orders did we sell products for less than 10 Euro?
- Which customers have more than one billing address?

# Solution – Non-relational (noSQL) databases (Column stores)

- Not using the relational model
- Running well on clusters
- Open-source
- Built for the 21st century web states
- Schemaless

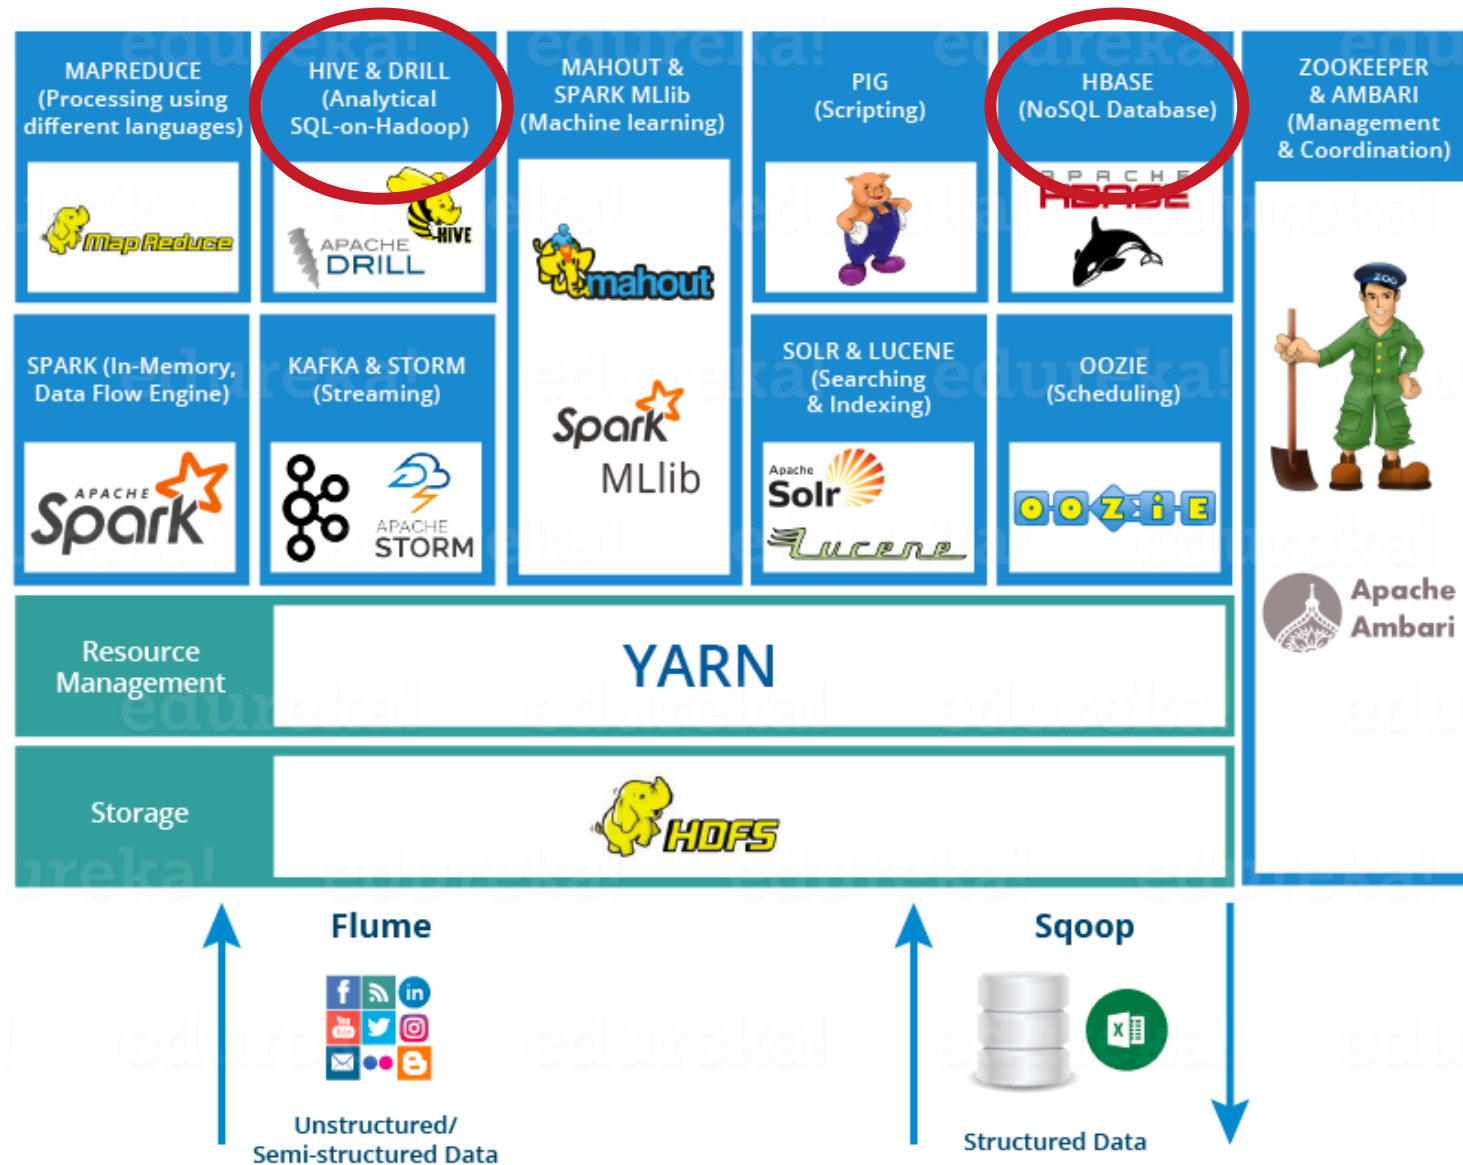
- Google and Amazon particularly influential in the rise of NoSQL databases
  - Google BigTable: powers Gmail, Google Maps, and other services
  - Amazon Dynamo: used e.g., by Amazon, Netflix
- Both capture huge amounts of data
- Both are running huge clusters of data
- Also popular: Open source solutions, e.g., Cassandra, HBase

# Some recent systems



Ali Davoudian, Liu Chen,  
and Mengchi Liu. 2018. A  
Survey on NoSQL Stores.  
*ACM Comput. Surv.* 51, 2,  
Article 40 (March 2019), 43  
pages.

# Hadoop ecosystem



- Relational databases (**Row** store):
  - Rows are stored consecutively
  - Optimal for row-wise access (e.g. SELECT \*)
- Non-relational databases (**Column** store):
  - Columns are stored consecutively
  - Optimal for attribute focused access (e.g. SUM, GROUP BY)

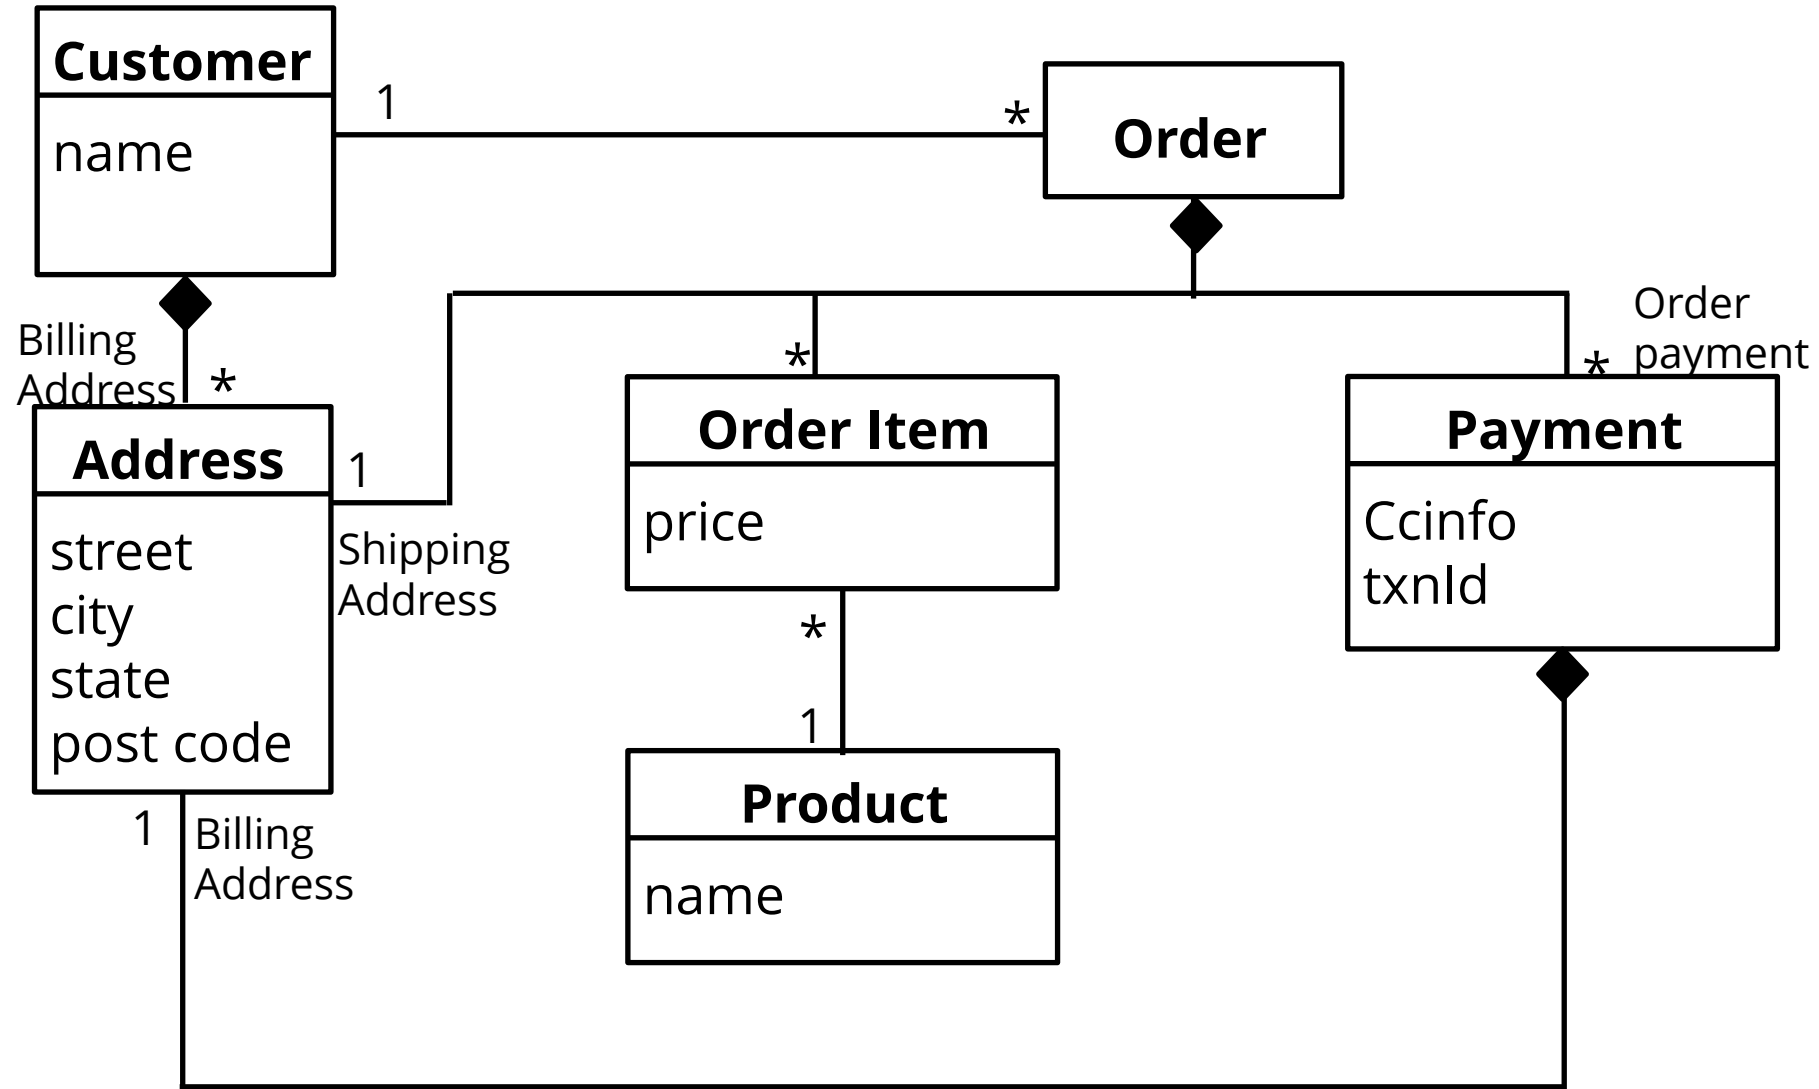
Paolo Atzeni, Francesca Bugiotti, Luca Cabibbo,  
Riccardo Torlone. Data Modeling in the NoSQL World.  
*Computer Standards and Interfaces*, 2020, 67, pp.103149.  
10.1016/j.csi.2016.10.003 . hal- 01611628



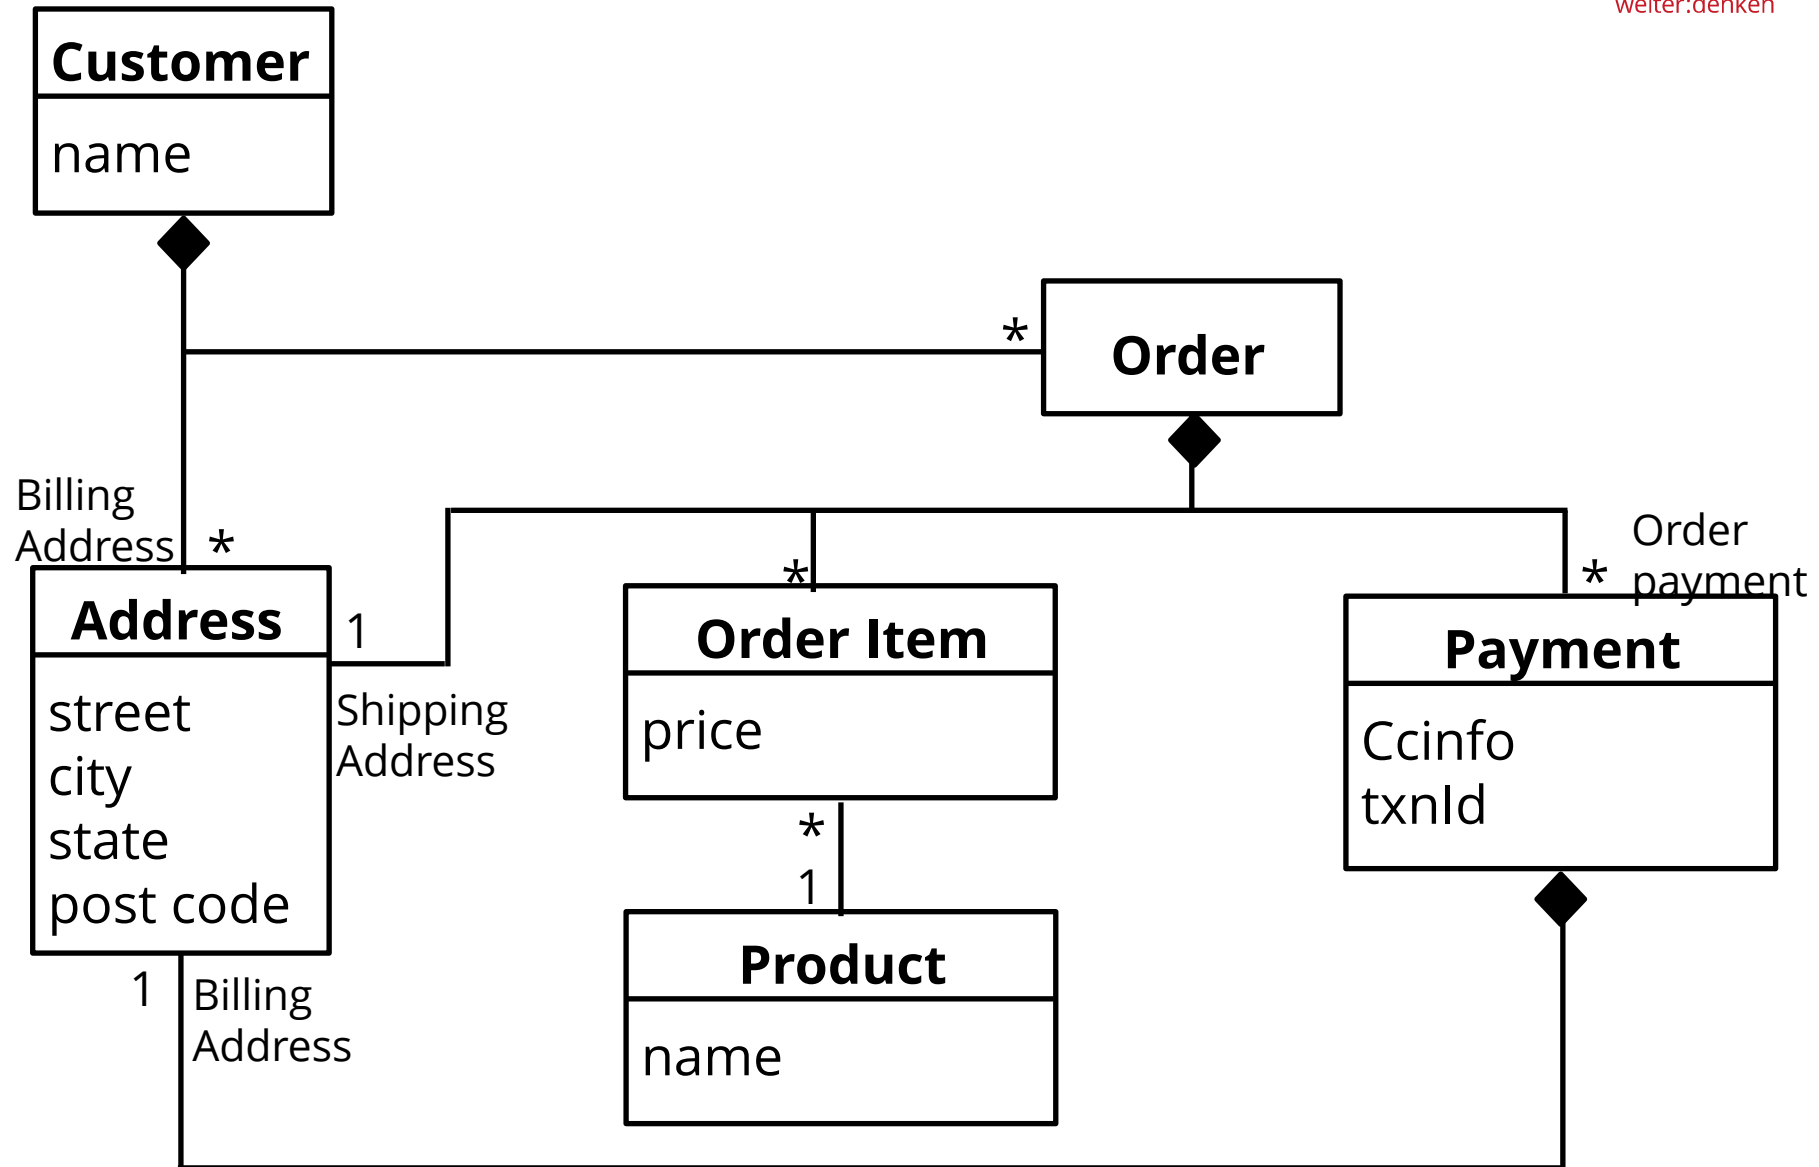
# Approach: Aggregates

- Data is **aggregated** in units (records) that have more complex design than tuples
- *How* to aggregate depends on the application
  - Sometimes internal validation mechanisms that decide whether aggregate is acceptable

# Example: Aggregate data model



# Example: Aggregate data model (Alternative)



- Aggregate-oriented databases work best when most data interaction is done with the same aggregate.
- An aggregate structure may help with some data interactions but be an obstacle for others.
- Aggregates make it easier for the database to manage data storage over clusters.

- **Relational Databases:**

- Allow you to manipulate any combination of rows from any table in a single transaction.
- ACID transactions: atomic, consistent Isolated, and Durable

- **Aggregate-oriented databases:**

- Support atomic manipulation of a single aggregate at a time
- No ACID transactions that span multiple aggregates
- We have to manage the atomic manipulation of multiple aggregates in our application code!

- Part 1:
  - Column stores (Introduction)
  - **Big Table**
  - HBase and Hive
- Part 2:
  - Job Scheduling
  - Coordination

- Lots of (semi-)structured data at Google
  - URLs:
    - Contents, crawl metadata, links, anchors, pagerank, ...
  - Per-user data:
    - User preference settings, recent queries/search results, ...
  - Geographic locations:
    - Physical entities (shops, restaurants, etc.), roads, satellite image data, user annotations, ...
- Scale is large
  - Billions of URLs, many versions/page (~20K/version)
  - Hundreds of millions of users, thousands or q/sec
  - 100TB+ of satellite image data

- Distributed multi-level map
- Fault-tolerant, persistent
- Scalable
  - Thousands of servers
  - Terabytes of in-memory data
  - Petabyte of disk-based data
  - Millions of reads/writes per second, efficient scans
- Self-managing
  - Servers can be added/removed dynamically
  - Servers adjust to load imbalance

*Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach, Mike Burrows, Tushar Chandra, Andrew Fikes, and Robert E. Gruber. 2008. Bigtable: A Distributed Storage System for Structured Data. ACM Trans. Comput. Syst. 26, 2, Article 4 (June 2008), 26 pages.*

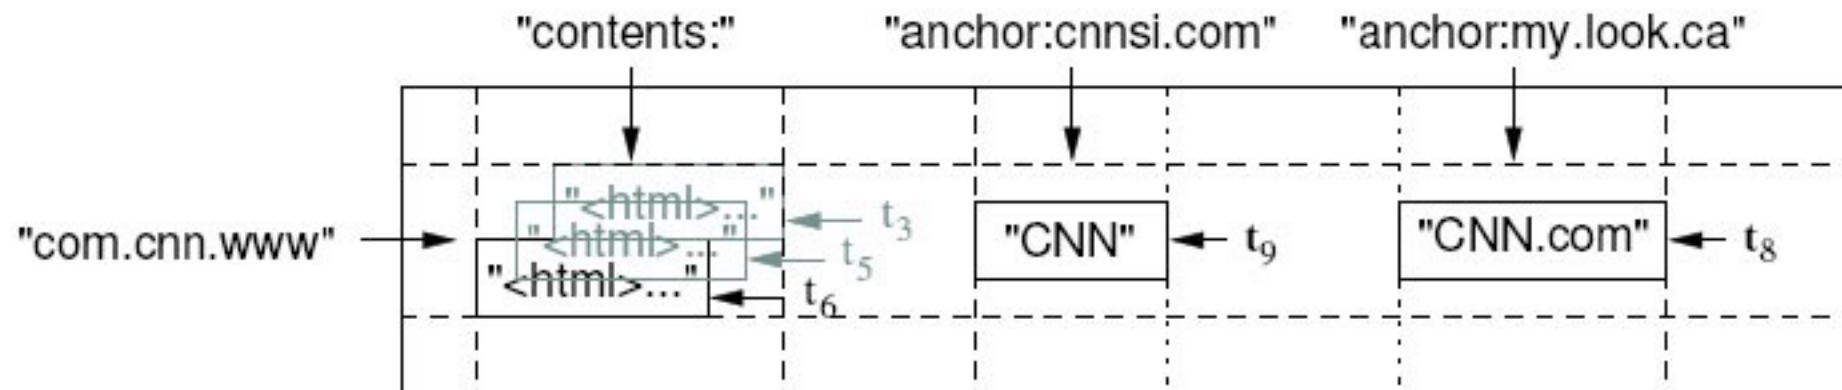


- Building blocks:
  - Google File System (GFS): Raw storage
  - Scheduler: schedules jobs onto machines
  - Lock service: distributed lock manager
  - MapReduce: simplified large-scale data processing
  
- BigTable's use of building blocks:
  - GFS: stores persistent data (SSTable file format for storage of data)
  - Scheduler: schedules jobs involved in BigTable serving
  - Lock service: master election, location bootstrapping
  - Map Reduce: often used to read/write BigTable data

# Basic Data Model


- A BigTable is a sparse, distributed persistent multi-dimensional sorted map

*(row, column, timestamp) -> cell contents*



- Good match for most Google applications

# Application of BigTable



- Part 1:
  - Column stores (Introduction)
  - Big Table
  - **HBase and Hive**
- Part 2:
  - Job Scheduling
  - Coordination

# History of HBase

- Started at the end of **2006**
- Modelled after **Google's Bigtable paper** (2006)
- **January 2008:** Hadoop becomes Apache top level project, HBase becomes subproject
- **May 2010:** HBase becomes an Apache top level project
- Contributors from Cloudera, Facebook, Intel, Hortonworks, etc.

Aiyer, Amitanand S., Mikhail Bautin, Guoqiang Jerry Chen, Pritam Damania, Prakash Khemani, Kannan Muthukkaruppan, Karthika Ranganathan, Nicolas Spiegelberg, Liyin Tang and Madhuwanti Vaidya. "Storage Infrastructure Behind Facebook Messages: Using HBase at Scale." *IEEE Data Eng. Bull.* 35 (2012): 4-13.

- “Apache HBase is the Hadoop database, a distributed, scalable, big data store.”
- On top of HDFS
- Billions of rows \* millions of columns
- Non-relational DB / NoSQL
- Column store: columns instead of tables
- Key-value cells

# How HBase works



- **Structured data**, scaling to petabytes
- **Efficient** handling of **diverse** data
  - wrt data size (URLs, web pages, satellite images)
  - wrt latency (backend bulk processing vs. real-time data serving)
- **Efficient** read and write of **individual records**



# HBase vs. Hadoop

- Hadoop's use case is **batch processing**
  - Not suitable for a single record lookup
  - Not suitable for adding **small amounts** of data at all times
  - Not suitable for making **updates** to existing records
- HBase addresses Hadoop's weaknesses
  - Provides fast lookup of **individual** records
  - Supports **insertion** of **single** records
  - Supports record **updating**
  - Not all columns are of interest to everyone; each client only wants a particular subset of columns (**column-based storage**)

# HBase vs. Hadoop

HBase is built on top of HDFS!

	Hadoop	HBase
writing	file append only, no updates	random write, updating
reading	sequential	random read, small range scan, full scan
structured storage	up to the user	sparse column family data model

# HBase vs. Hadoop

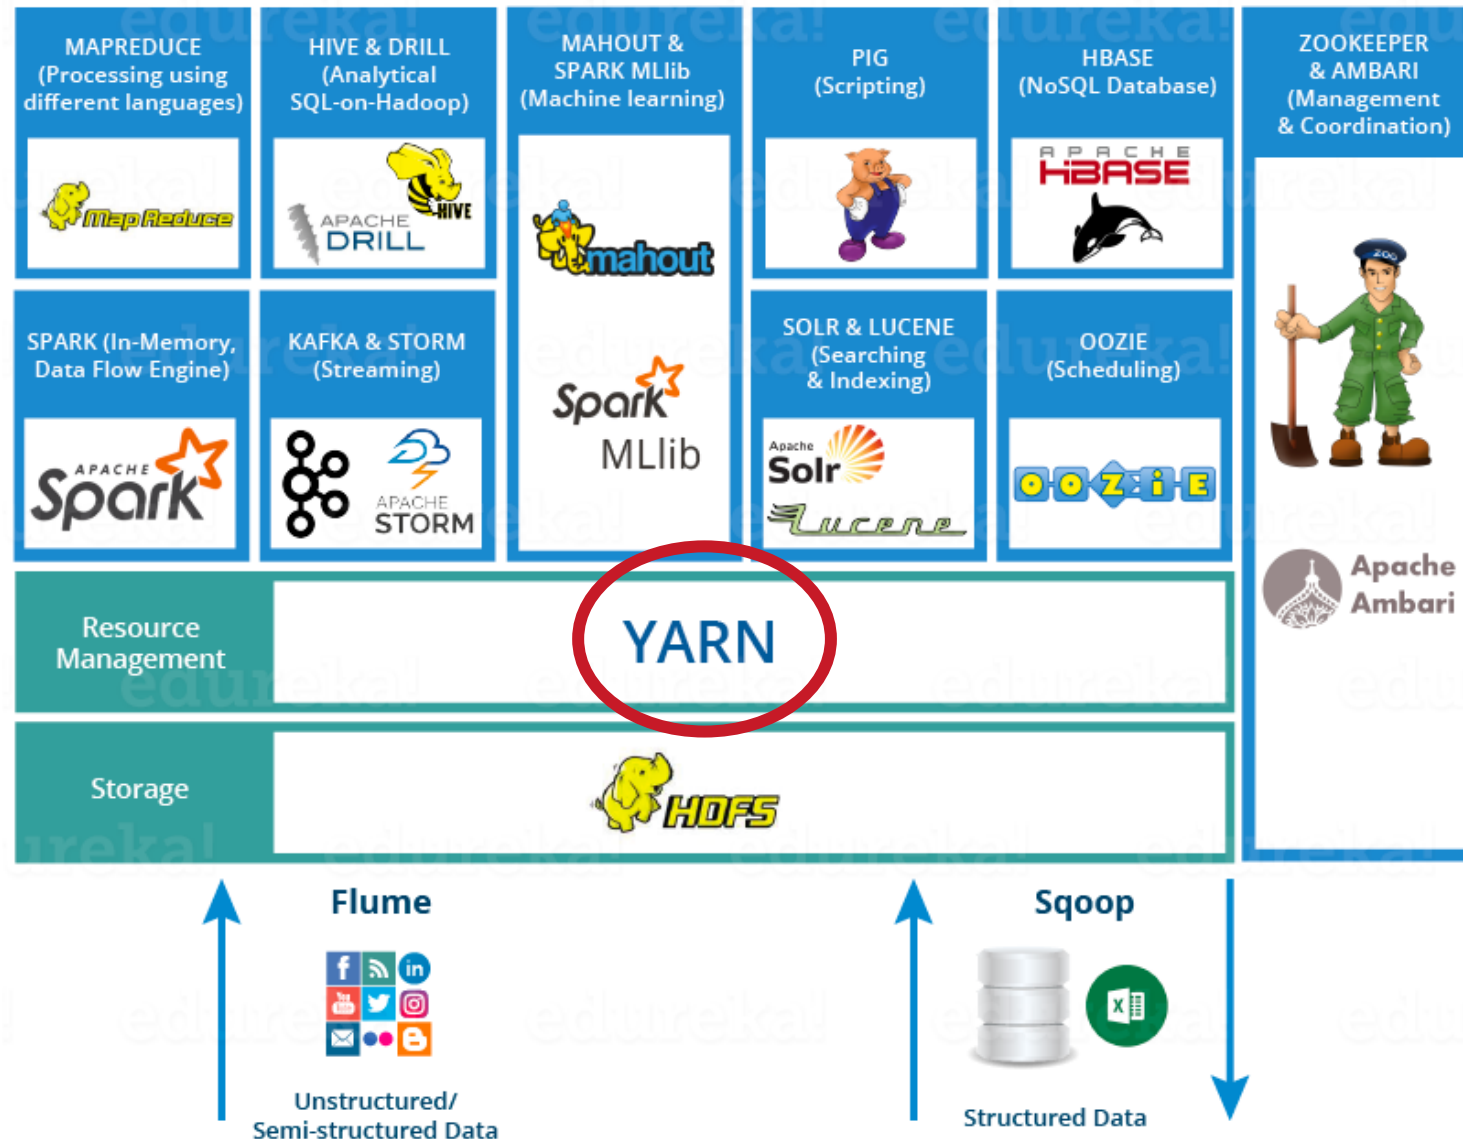
	small to medium-volume applications	use when scaling up in terms of dataset size, read/write concurrency
	RDBMS	HBase
schema	fixed	random write, updating
orientation	row-oriented	column-oriented
query language	SQL	simple data access model
size	terabytes (at most)	billions of rows, millions of columns
scaling up	difficult (workarounds)	add nodes to a cluster

- Developed at Facebook
- Used for majority of Facebook jobs
- “Relational database” built on Hadoop
  - Maintains list of table schemas
  - SQL-like query language (HiveQL)
  - Can call Hadoop Streaming scripts from HiveQL
  - Supports table partitioning, clustering, complex data types, some optimizations


- Part 1:
  - Column stores (Introduction)
  - Big Table
  - HBase and Hive
- **Part 2:**
  - **Job Scheduling**
  - Coordination

- A Hadoop job consists of Map tasks and Reduce tasks
- Multiple customers with multiple jobs
  - Users/jobs = “tenants”
  - Multi-tenant system
- Need a way to schedule all these jobs (and their constituent tasks)
- Need to be *fair* across the different tenants

# Hadoop ecosystem



# Scheduling




*Which tasks run when?*

Processor

Task	Length	Arrival
1	10	0
2	5	6
3	3	8



# FIFO Scheduling (First-In First-Out) / FCFS




- *Maintain tasks in a queue in order of arrival*
- *When processor free, dequeue head and schedule it*

# FIFO/FCFS Performance

- Average completion time may be high
- For our example on previous slides,
  - average completion time of FIFO/FCFS =  
 $(\text{Task 1} + \text{Task 2} + \text{Task 3})/3$   
 $= (10+15+18)/3$   
 $= 43/3$   
 $= 14.33$

# STF Scheduling (Shortest Task First)




- Maintain all tasks in a queue, in increasing order of running time
- When processor free, dequeue head and schedule

# STF is Optimal!

- Average completion of STF is the shortest among all scheduling approaches!
- For our example on previous slides,
  - Average completion time of STF =
    - $(\text{Task 1} + \text{Task 2} + \text{Task 3})/3$
    - $= (18+8+3)/3$
    - $= 29/3$
    - $= 9.66$
    - (versus 14.33 for FIFO/FCFS)
- In general, STF is a special case of priority scheduling
  - Instead of using time as priority, scheduler could use user-provided priority

# Round-Robin Scheduling




- Use a quantum (say 1 time unit) to run portion of task at queue head
- Pre-empt processes by saving their state, and resuming later
- After pre-empting, add to end of queue

# Round-Robin vs. STF/FIFO

- Round-Robin preferable for
  - Interactive applications
  - User needs quick responses from system
- FIFO/STF preferable for Batch applications
  - User submits jobs, goes away, comes back to get result

- Provides resource management to support processing petabytes of data.
- Hadoop YARN has two popular schedulers
  - *Hadoop Capacity Scheduler*
  - *Hadoop Fair Scheduler*



Murthy et al. Apache Hadoop YARN: Moving beyond MapReduce and Batch Processing with Apache Hadoop 2, Addison-Wesley, 2014

# Hadoop Capacity Scheduler

- Contains multiple queues
- Each queue contains multiple jobs
- Each queue guaranteed some portion of the cluster capacity, e.g.,
  - Queue 1 is given 80% of cluster
  - Queue 2 is given 20% of cluster
  - Higher-priority jobs go to Queue 1
- For jobs within same queue, FIFO typically used
- Administrators can configure queues



- Administrators can configure each queue with limits
  - Soft limit: how much % of cluster is the queue guaranteed to occupy
  - (Optional) Hard limit: max % of cluster given to the queue
- Elasticity
  - A queue allowed to occupy more of cluster if resources free
  - But if other queues below their capacity limit, now get full, need to give these other queues resources
- Pre-emption not allowed!
  - Cannot stop a task part-way through
  - When reducing % cluster to a queue, wait until some tasks of that queue have finished

- Queues can be hierarchical
  - May contain child sub-queues, which may contain child sub-queues, and so on
  - Child sub-queues can share resources equally
- Scheduling can take memory requirements into account (memory specified by user)

# Hadoop Fair Scheduler

- Goal: all jobs get equal share of resources
- When only one job present, occupies entire cluster
- As other jobs arrive, each job given equal % of cluster
  - E.g., each job might be given equal number of cluster-wide YARN containers
  - Each container == 1 task of job

# Hadoop Fair Scheduler (2)

- Divides cluster into pools
  - Typically one pool per user
- Resources divided equally among pools
  - Gives each user fair share of cluster
- Within each pool, can use either
  - Fair share scheduling, or
  - FIFO/FCFS
  - (Configurable)

- Some pools may have *minimum shares*
  - Minimum % of cluster that pool is guaranteed
- When minimum share not met in a pool, for a while
  - Take resources away from other pools
  - By pre-empting jobs in those other pools
  - By *killing* the currently-running tasks of those jobs
    - Tasks can be re-started later
    - Ok since tasks are idempotent!
  - To kill, scheduler picks most-recently-started tasks
    - Minimizes wasted work

# Other HFS Features

- Can also set limits on
  - Number of concurrent jobs per user
  - Number of concurrent jobs per pool
  - Number of concurrent tasks per pool
- Prevents cluster from being hogged by one user/job

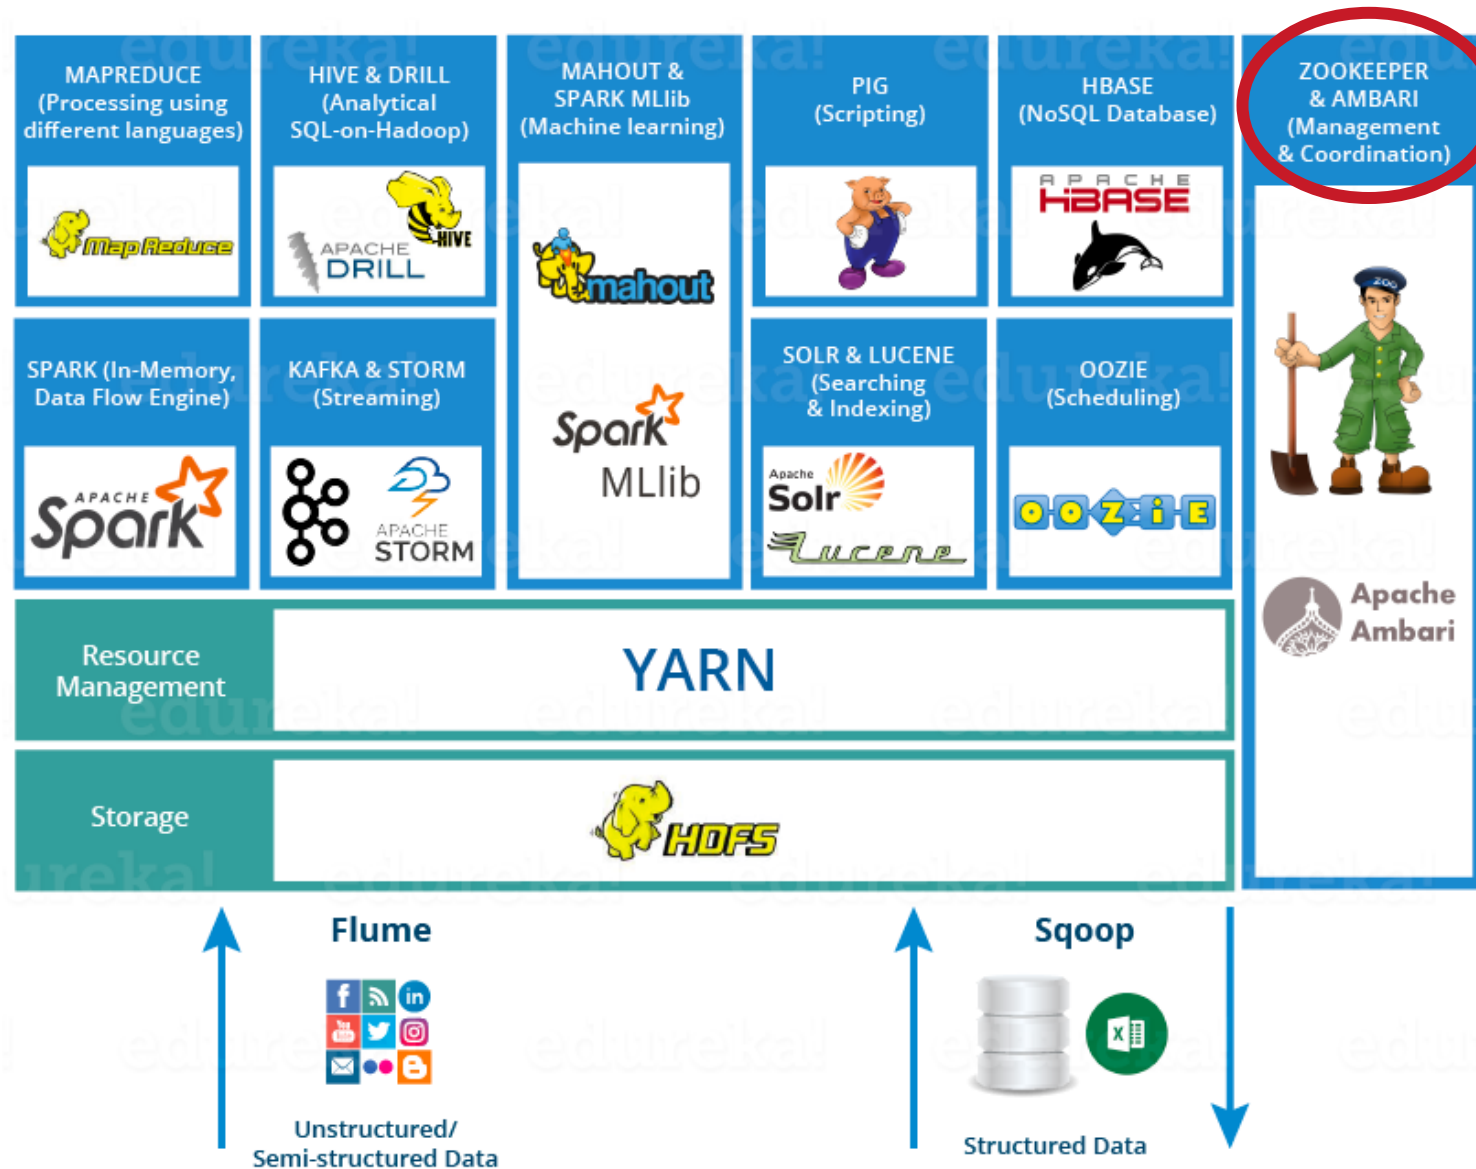
- Part 1:
  - Column stores (Introduction)
  - Big Table
  - HBase and Hive
- **Part 2:**
  - Job Scheduling
  - **Coordination**

# Challenge

- In the past: a **single** program running on a **single** computer with a **single** CPU
- Today: applications consist of **independent** programs running on a **changing** set of computers
- Difficulty: **coordination** of those independent programs
- Developers have to deal with **coordination logic** and **application logic** at the same time



# Hadoop ecosystem



# Apache Zookeeper

- ZooKeeper is a centralized service for maintaining configuration information, naming, providing distributed synchronization, and providing group services.
- <https://zookeeper.apache.org/>




# What is Zookeeper?



- ZooKeeper service is an ensemble of servers that use replication (high availability)
- (meta)data is cached on the client side
  - ID of datanodes, instead of probing ZooKeeper every time.
  - What if data changes?
    - Polling?
    - Watch mechanism: clients can watch for an update of a given object

- ZooKeeper data is replicated on each server that composes the service
  - Recovery from master failures
- Clients connect to exactly one server to submit requests
  - **read** requests served from the local replica
  - **write** requests are processed by an agreement protocol (an elected server leader initiates processing of the write request)

# Zookeeper Servers



- All servers store a copy of the data (in memory)
- A leader is elected at startup
- Followers service clients, all updates go through leader
- Update responses are sent when a majority of servers have persisted the change

- Big Table
- HBase and Hive
- Job Scheduling
- Coordination

# What's next – Distributed Data Processing

- Map/Reduce framework
- Querying
  - Spark Core API
  - Pig, Pig Latin
- Machine Learning at Scale