

➤ Big Data

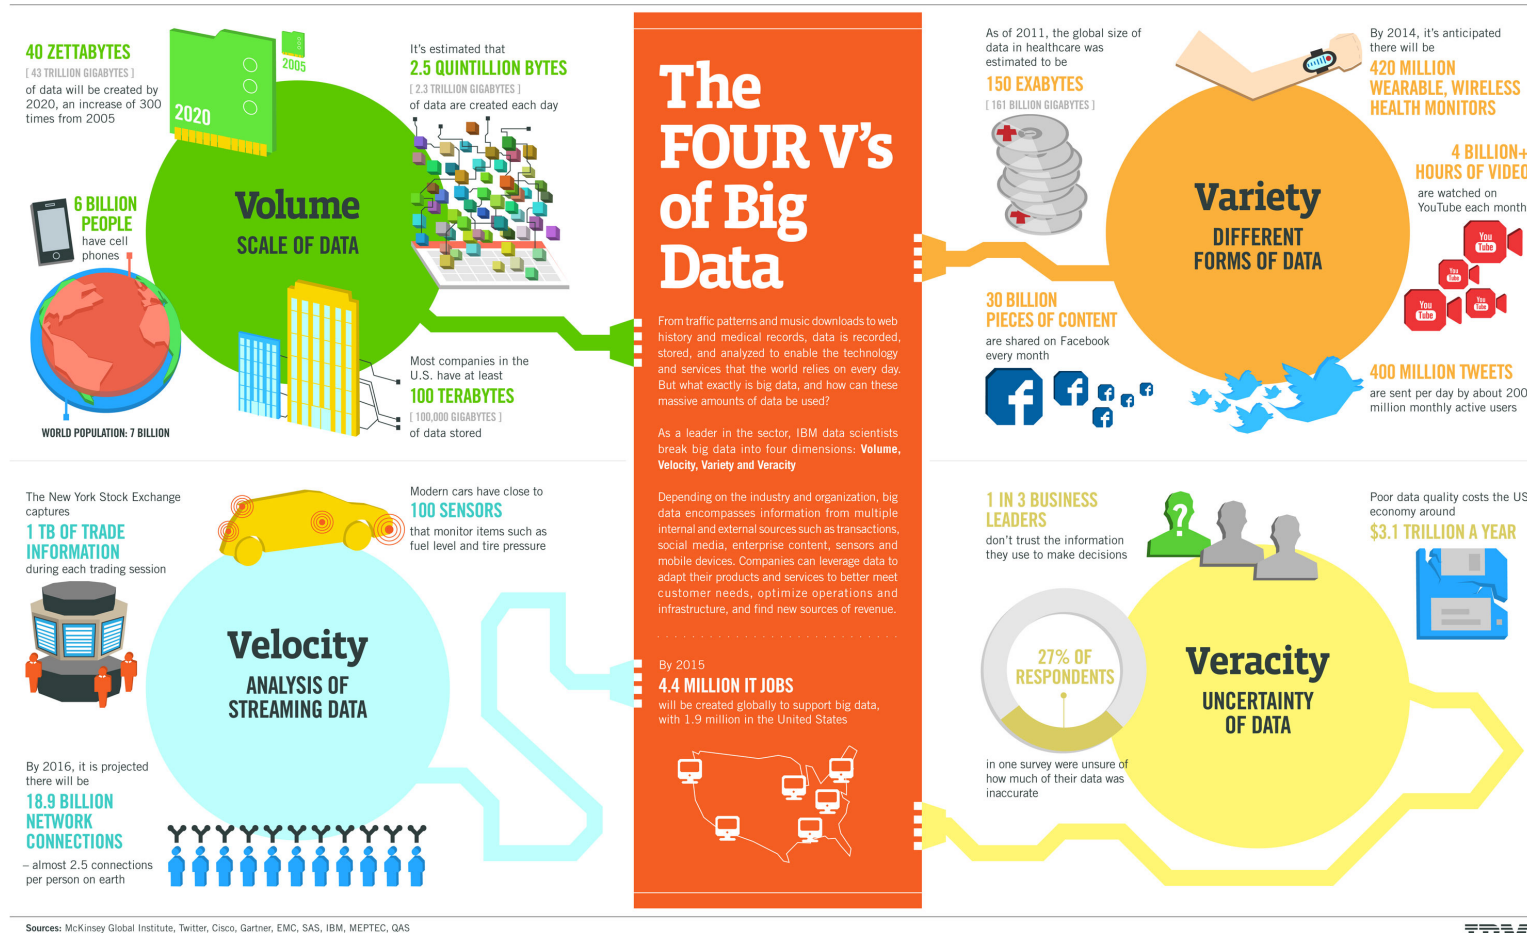
Session 5: Processing Large Data Streams

Frank Hopfgartner
Institute for Web Science and Technologies

Last week

- Map/Reduce framework
- Querying
 - Spark Core API
 - Pig, Pig Latin
- Machine Learning at Scale
 - Spark MLlib
 - Mahout

Recap: Four V's of Big Data



Motivation

- So far we have really just talked about processing historical, existing big data
 - Sitting on HDFS
 - Sitting in a database
- But how does new data get into your cluster? Especially if it is 'big data'?
- Streaming lets you publish this data, in real-time, to your cluster
 - And you can even process it in real-time as it comes in.

Intended Learning Outcomes


At the end of this lecture, you will be able to:

- Outline use cases to stream data
- Explain record-at-a-time streaming
- Distinguish between various declarative streaming cases

- Introduction
- Record-at-a-time streaming
- Declarative, functional streaming
- Declarative, relational streaming

What is streaming?


Continuously integrating new, infinitely large data to compute results



Processing steps in pipeline fashion

Given a stream source

- The amount of data received is unlimited in size
- The volume of data is continuous/variable
 - Typically unpredictable
- Example: Twitter



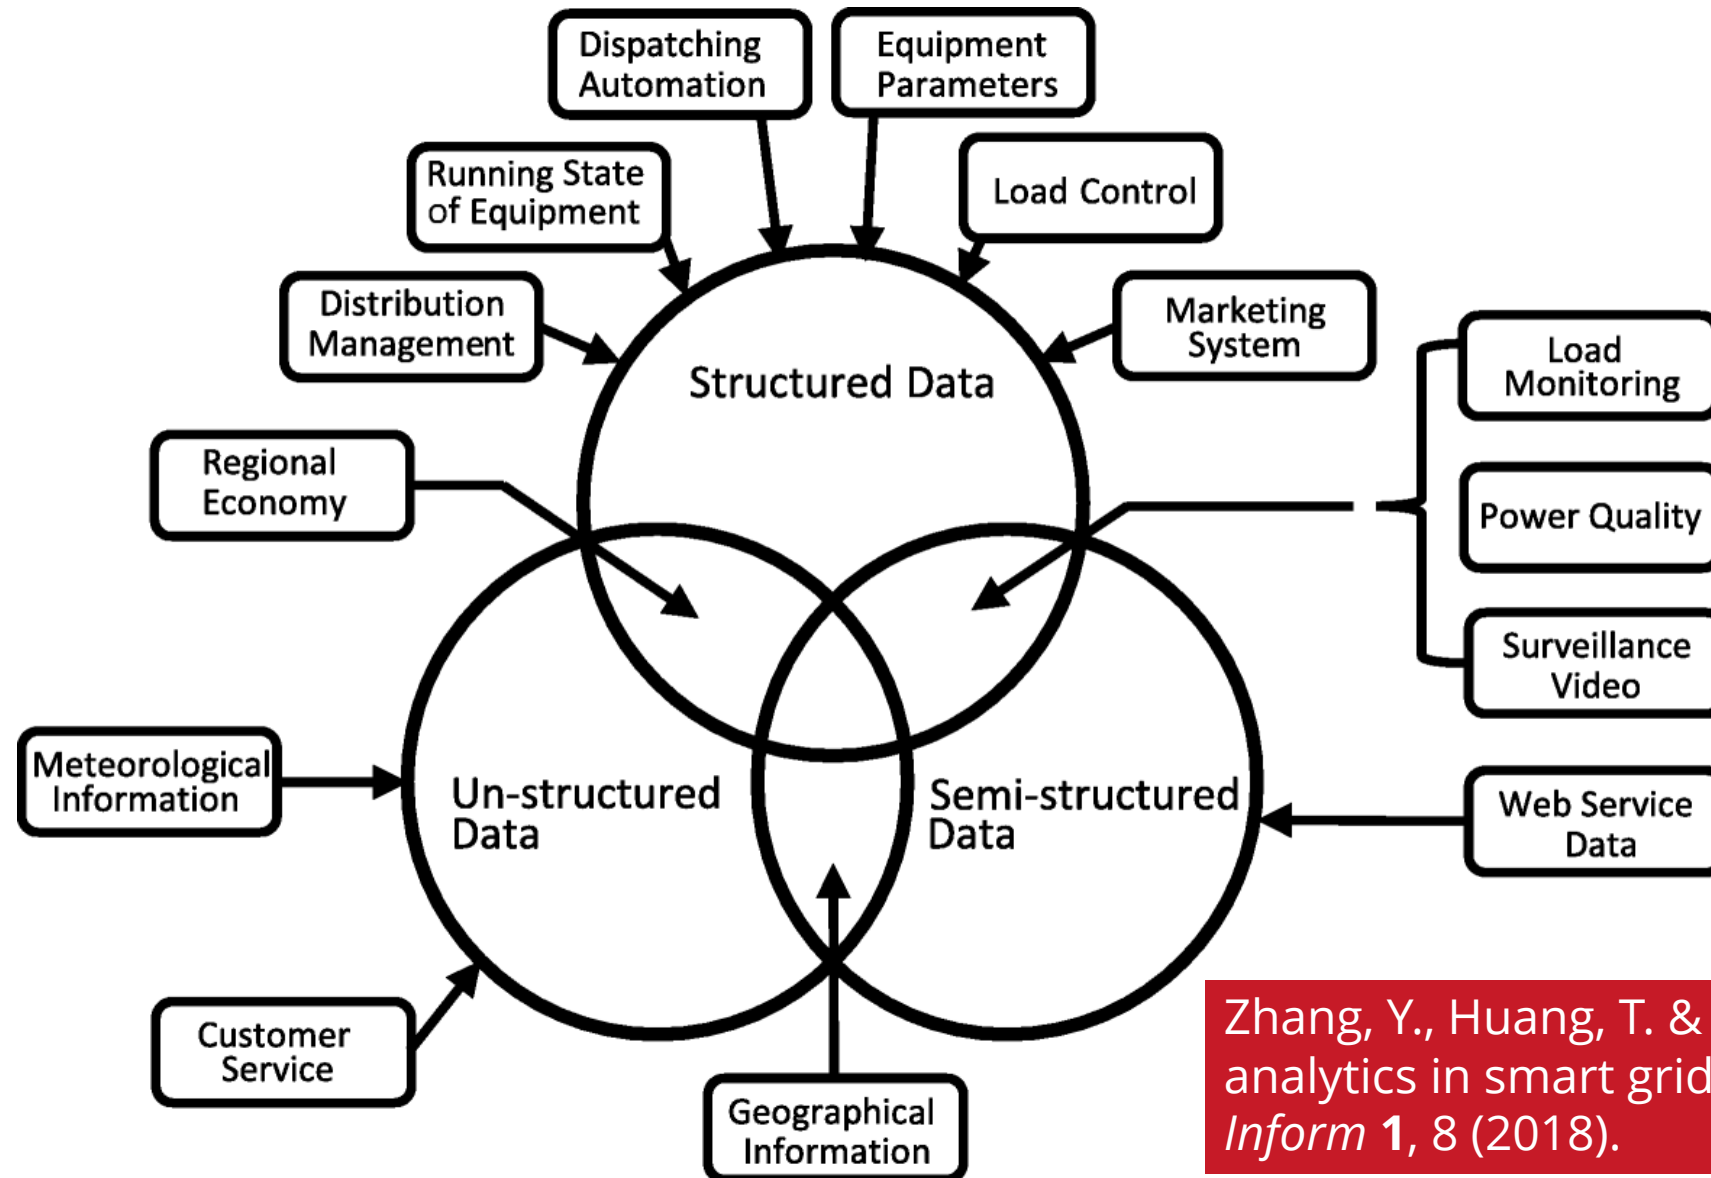
Stream processing use cases

1. Notifications and alerting
 - Redundant failures to log in
2. Real-time reporting
 - Dashboards, e.g. for production/traffic flow, system load, uptime,....
3. Incremental ETL
 - Update the data warehouse
 - Maintain correctness (do not lose data or add data twice)
4. Update data to serve in real time
 - e.g. updating a key-value or relational store with statistics
5. Real-time decision making
 - Analyzing new inputs, e.g. analyzing credit card transactions to discover and prevent fraud
6. Online machine learning
 - e.g. learning to recommend on changing platforms

Example: Thailand's Tsunami warning system



Example: Smart Grids



Zhang, Y., Huang, T. & Bompard, E.F. Big data analytics in smart grids: a review. *Energy Inform* 1, 8 (2018).

Advantages of streaming

- Lower latency
 - Timescale: minutes, seconds, milliseconds
 - No need to re-process all previously occurred data
 - Retaining **state of system**
- Higher efficiency for updates
 - Compared to repeated batch jobs

Batch vs. Stream: Same, same, but different

Same functionality needed

- Computing account balances
- Computing statistics
- ...

Batch processing

- Fixed data during processing
- Query triggers processing

Stream processing

- Continuous arrival of new data
- New data triggers processing (mostly)

Challenges of stream processing

- Batch processing is simpler to understand, troubleshoot and program
- Batch processing allows for much higher throughput
- Data may arrive out of order, e.g.,

```
{value: 1, time: "2017-04-07T00:00:00"}  
{value: 2, time: "2017-04-07T01:00:00"}  
{value: 5, time: "2017-04-07T02:00:00"}  
{value: 10, time: "2017-04-07T01:30:00"}  
{value: 7, time: "2017-04-07T03:00:00"}
```

- Consider a question like:
 - Did 2 – 10 – 5 appear?
 - Did 2 – 10 – 5 **not** appear?
- Remember events? For how long?
- Remember states?

Challenges of stream processing (2)

- Processing out-of-order data
based on application timestamps (*event time*)
- Maintaining large amounts of state
- Supporting high-data throughput
- Processing each event exactly once despite machine failures
- Handling load imbalance and stragglers
- Responding to events at low latency
- Joining with external data in other storage systems
- Determining how to update output sinks as new events arrive
- Writing data transactionally to output systems
- Updating your application's business logic at runtime

Record-at-a-Time:

- API just hands over one record-at-a-time to application
- Application handles all challenges
- Apache Storm

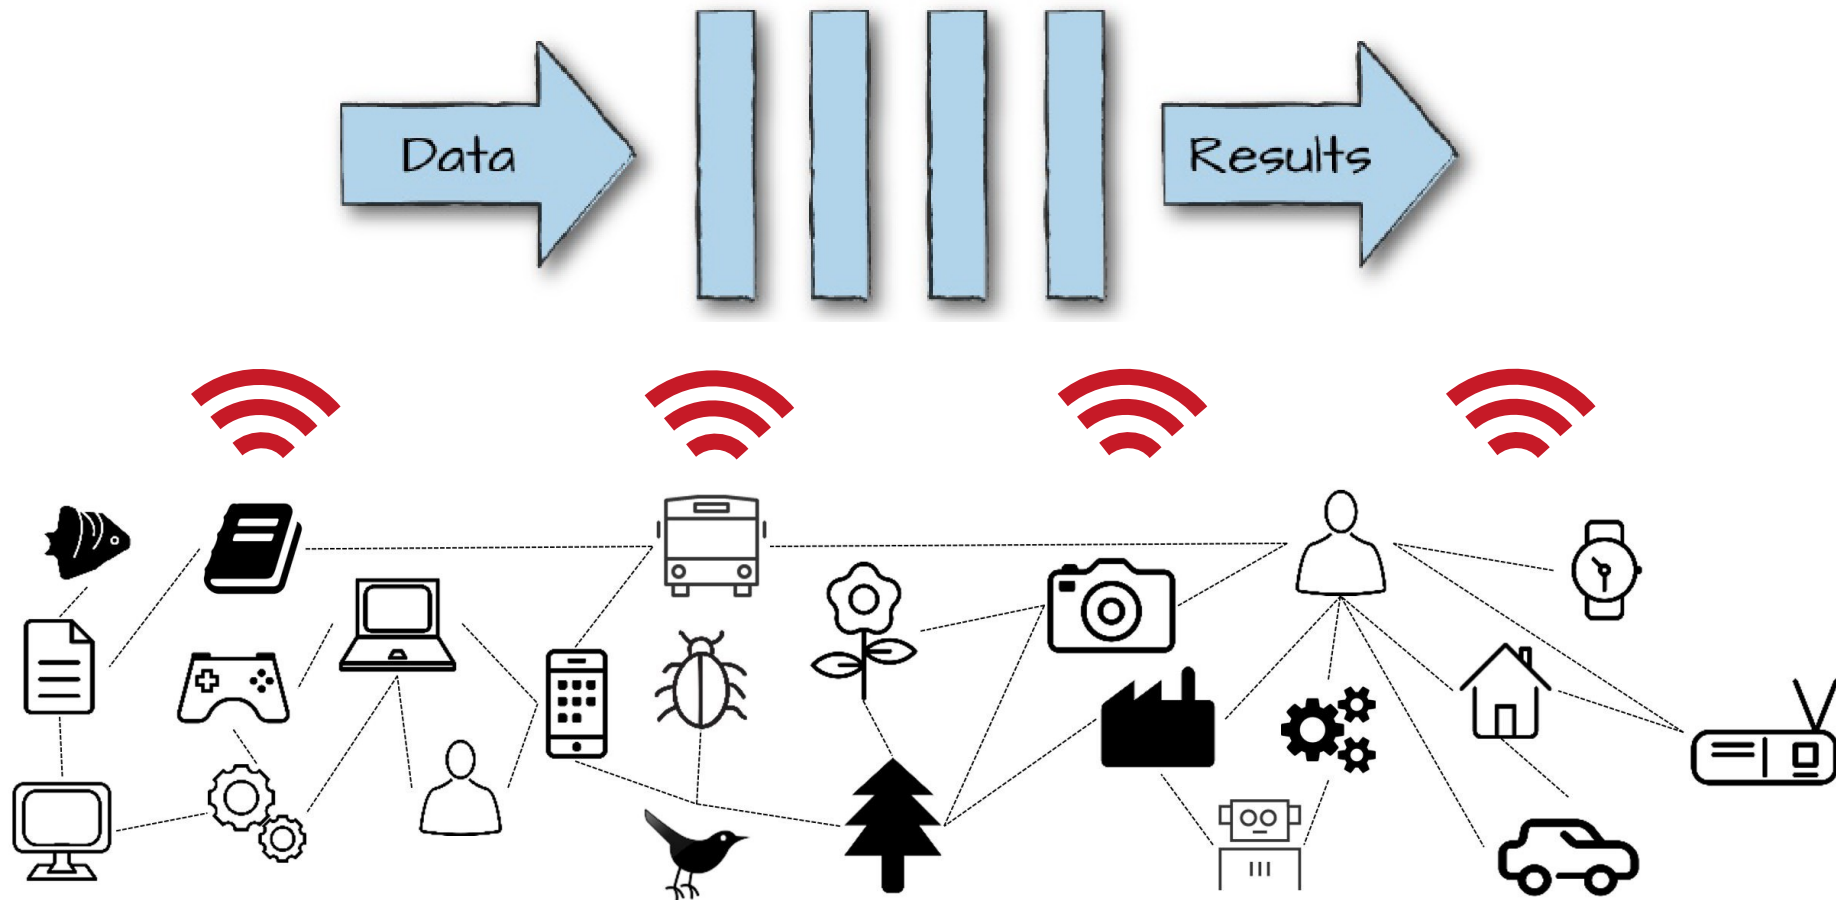
Declarative, functional API:

- Describe what to compute, not how
- Functional: *map, reduce, filter*
- Dstreams API, Google Dataflow, Apache Kafka

Declarative, relational API:


- Rich automatic optimization of execution (beyond functional)
- Spark Structured Streaming, Apache Flink

Event time vs. processing time




Web of Things

Continuous vs. batch processing



One record at a time



Microbatches of DataFrames

Continuous processing

- Processes data immediately
- Lower base latency

Microbatch processing

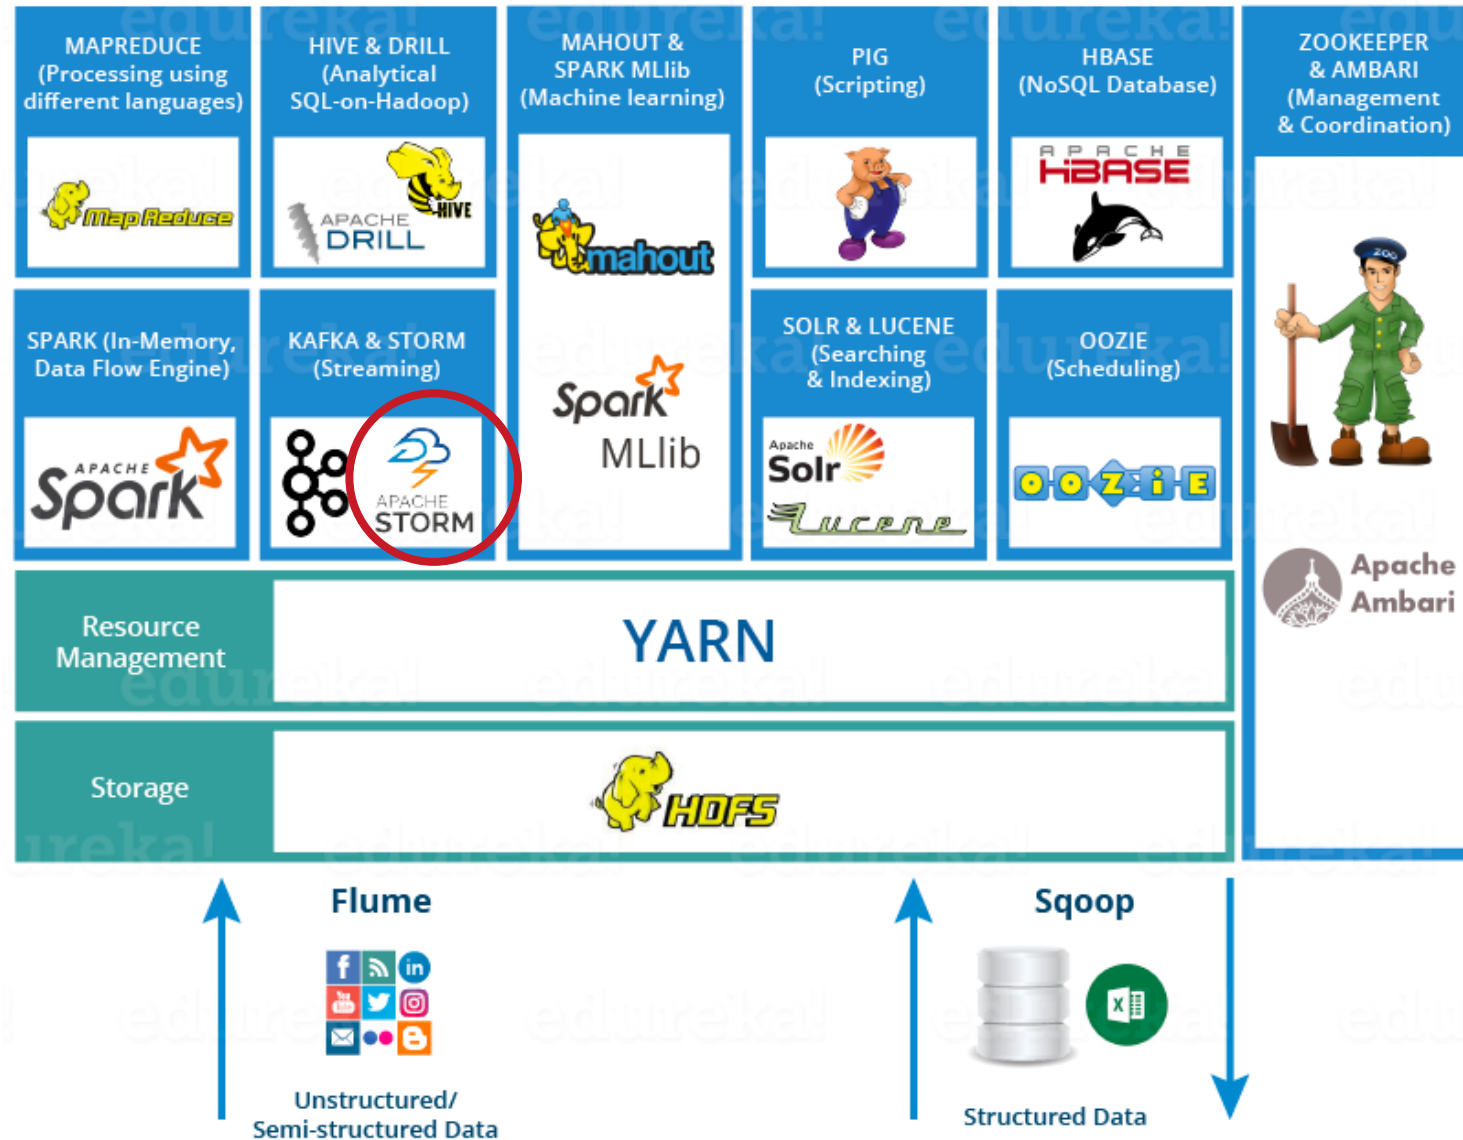
- Waits for some amount of data to arrive before processing
- Latencies starting at 100ms to 1s
- Improved throughput
- Preferred by distributed streaming
 - If scalability is an issue, throughput must be optimized

Practical, continuous applications

- reacts to data in real time
- mixes
 - Streaming jobs
 - Batch jobs
 - Joins between streaming and offline data
 - Interactive ad-hoc queries

- Introduction
- **Record-at-a-time streaming**
 - **Apache Storm**
- Declarative, functional streaming
- Declarative, relational streaming

Hadoop ecosystem



Apache Storm


- Distributed and fault-tolerant real-time computation system for processing limitless streaming data
- Built at Twitter
- Real-time analytics
 - Not batch data processing like Hadoop
- Define a topology: graph of computation
 - Consumes streams of data and processes those streams in arbitrarily complex ways, repartitioning the streams as needed



Storm vs Hadoop


Storm	Hadoop
Real-time streams of data	Batch data processing
Stateless	Stateful (data stored on HDFS)
Zookeeper coordination	Zookeeper coordination
1K msg / sec processed	TB/PB processed in minutes/hours
Topology runs as more data arrives	M/R jobs completed and results written on HDFS

- Two kinds of nodes: spouts and bolts
 - **Spout**: source of data streams
 - **Bolt**: process input stream and outputs new stream
- Nodes execute in parallel



- **Spout**
 - Data sources like Twitter Streaming API
 - Kafka queue (see later)
 - Read from datasources
- **Bolt**
 - Filtering, aggregation, joining operations
 - Interact with other datasources, e.g., databases
- **Topology**
 - Directed graph where vertices are computation and edges are stream of data
 - Distributed over multiple worker nodes running all the time and waiting for jobs to process
 - Multiple nodes can execute one bolt and take a share of the data
 - Topology is run by the master node (called Nimbus) assigning tasks to nodes

Example: Twitter word count



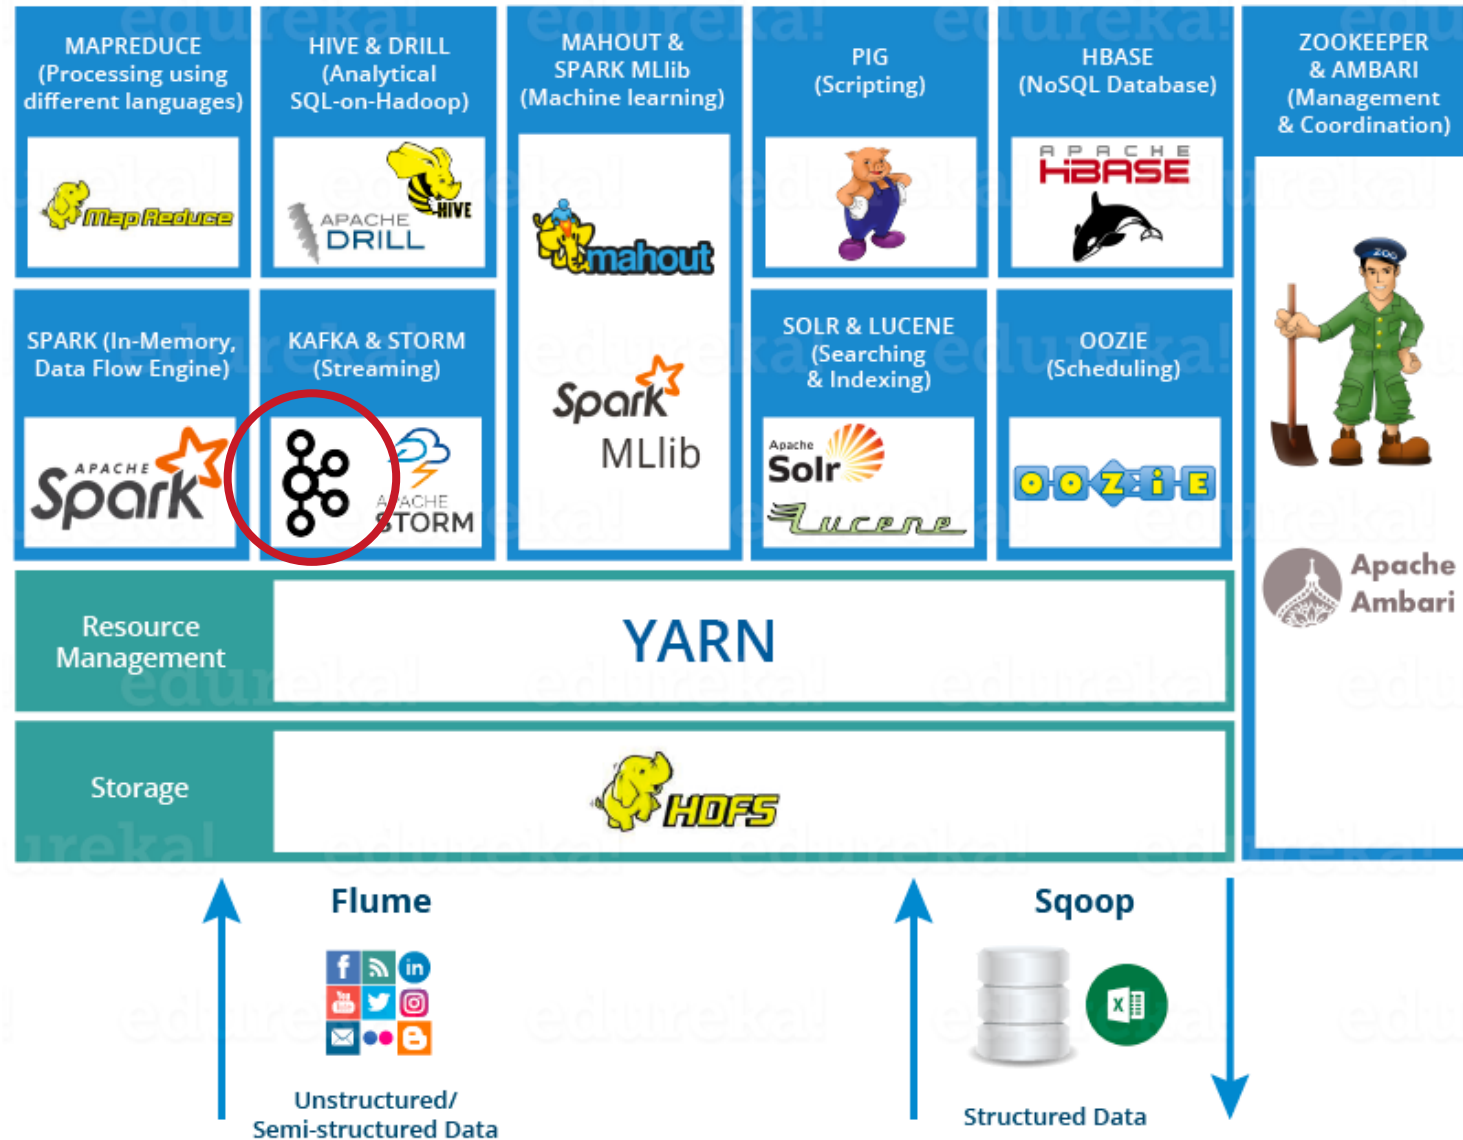
Storm – Use cases

	“Prevent” Use Cases	“Optimize” Use Cases
Financial Services	<ul style="list-style-type: none">✓ Securities fraud✓ Operational risks & compliance violations	<ul style="list-style-type: none">✓ Order routing✓ Pricing
Telecom	<ul style="list-style-type: none">✓ Security breaches✓ Network outages	<ul style="list-style-type: none">✓ Bandwidth allocation✓ Customer service
Retail	<ul style="list-style-type: none">✓ Shrinkage✓ Stock outs	<ul style="list-style-type: none">✓ Offers✓ Pricing
Manufacturing	<ul style="list-style-type: none">✓ Preventative maintenance✓ Quality assurance	<ul style="list-style-type: none">✓ Supply chain optimization✓ Reduced plant downtime
Transportation	<ul style="list-style-type: none">✓ Driver monitoring✓ Predictive maintenance	<ul style="list-style-type: none">✓ Routes✓ Pricing

<https://hortonworks.com/apache/storm/>

- Introduction
- Record-at-a-time streaming
- **Declarative, functional streaming**
 - **Apache Kafka**
 - **DStreams API**
- Declarative, relational streaming

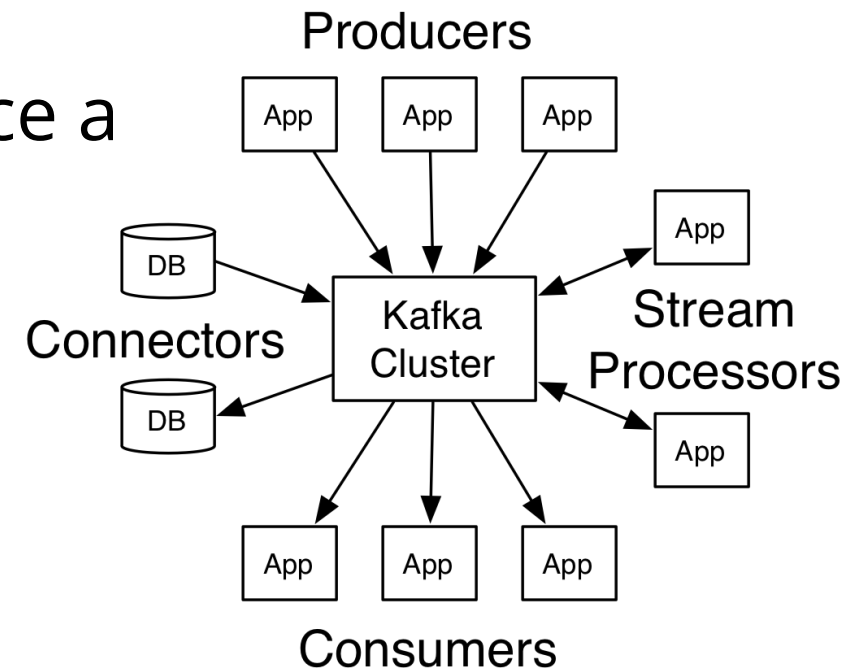
Hadoop ecosystem



- Originally built by LinkedIn (open-sourced in 2011)
- Kafka is a general-purpose **publish/subscribe messaging system**
- Kafka servers store all incoming messages from *publishers* for some period of time, and publishes them to a stream of data called a *topic*.
- Kafka *consumers* subscribe to one or more topics, and receive data as it is published
- A stream/topic can have many different consumers, all with their own position in the stream maintained
- It's not just for Hadoop.

Publish/subscribe model

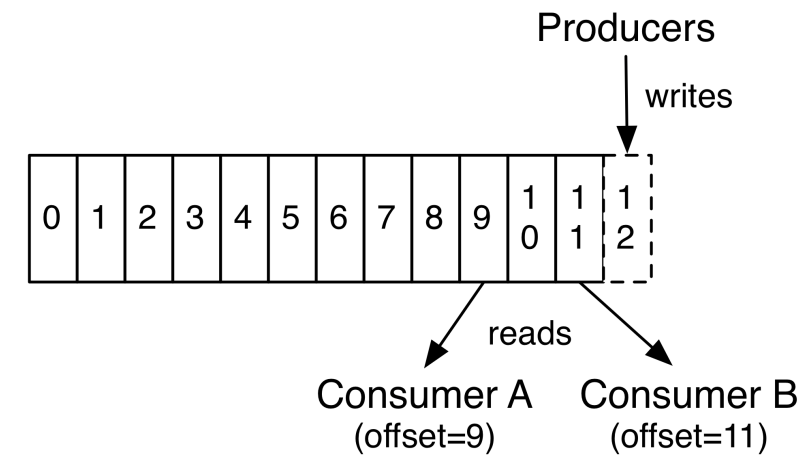
- Broadcast data to multiple processes
- Producers: publish a stream of data
- Consumers: subscribe to a stream
- Stream processors: consume and produce a new stream



How Kafka operates

- Kafka itself may be distributed among many processes on many servers
 - Will distribute the storage of stream data as well
- Consumers may also be distributed
 - Consumers of the same group will have messages distributed amongst themselves
 - Consumers of different groups will get their own copy of each message

- Producers write into a sequence of records that is continually appended
 - Sequences are partitioned and distributed in the cluster to scale
 - Partitions are replicated for fault tolerance
 - Order only maintained within a partition!
- Kafka cluster retains all published records for a predefined retention period (e.g., 2 days)
- Consumers read content using offset information



Example Kafka applications

- A retail application
 - takes in input streams of **sales** and **shipment** data
 - outputs a stream of **reorders** and **price adjustments** based on this data
- Usage at LinkedIn
 - <https://engineering.linkedin.com/kafka/kafka-linked-in-current-and-future>
 - Page views, clicks
 - <https://engineering.linkedin.com/kafka/running-kafka-scale>
 - Multiple datacenters
 - Mirroring data across Kafka clusters

Batch APIs

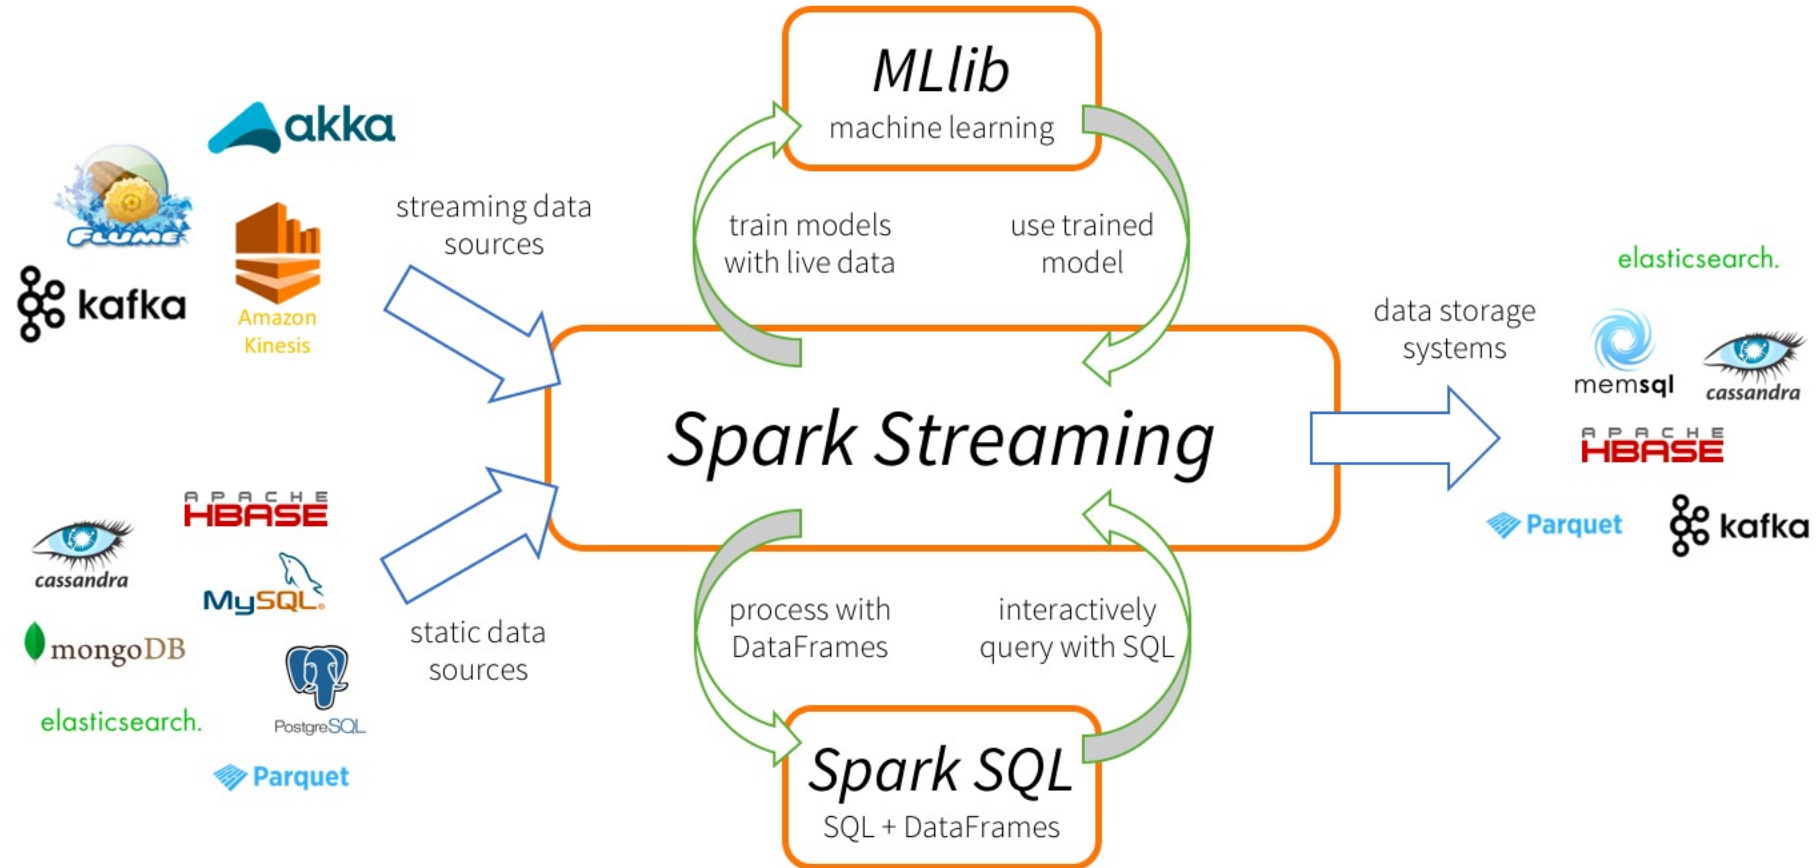
- RDD
 - Low level
- Data Frame
 - High level
 - Several means for self-optimization available

Streaming APIs

- Spark Streaming / DStreams API
 - Low level
- Structured Streaming API
 - High level
 - Several means for self-optimization available

- An extension of the core Spark
- Process real-time data from different data sources
- Stream of data divided into small batches (Discretized Stream – DStream)
 - Build on RDDs
- Can seamlessly integrate with other Spark components
- Scalable, high-throughput, fault-tolerant stream processing of live data streams

Spark streaming ecosystem



- Kafka
 - Most popular (and older)
 - High throughput (20k msg/sec)
- Apache Flume
 - Streaming event log data (web page visits, clicks) from a web server
 - Distributed / high availability
- Kinesis
 - Amazon AWS solution
- Streams are represented as a sequence of **RDDs**.


Spark – Stream processing

- Series of batch computations on small time intervals (windows over the stream)
- Spark Streaming receives live input data streams
- Divides the data into batches
- Spark engine processes batches




Discretised Streams (DStreams)

- Continuous stream of data
 - From source
 - Transforming an input file
- DStream is represented by a continuous series of RDDs
 - Each RDD has data from a certain interval
- Resilient Distributed Datasets (RDDs)
 - Keep data in memory
 - Can recover it without replication (track the lineage graph of operations that were used to build it)



Window operations

- Apply **transformations** (map, flatMap, etc) over a sliding window of data
- RDDs that fall within the window are combined and operated upon
 - Parameters: *window length, sliding interval*
 - Custom window-based transformations



- Streams arrive 24/7
- Storage able to recover from failures (HDFS)
 - Store computation metadata
 - Store data from streams
- When a node fails, each node in the cluster works to recompute part of the lost node's RDDs
- Batch interval needs to be set such that the expected data rate in production can be sustained

Spark Streaming – Use Cases

- Uber
 - Data from mobile users
 - Kafka as data source
 - Event data to structured data into HDFS
 - Analytics as M/R


- Pinterest
 - Real-time user interaction analysis
 - Use this for recommendations (products to buy, places to visit)

- Introduction
- Record-at-a-time streaming
- Declarative, functional streaming
- **Declarative, relational streaming**
 - **Spark Structured Streaming API**
 - **Apache Flink**

Spark Structured Streaming

- More recent
- More types of optimization (compared to Dstreams)
- Native support for event time
- Microbatches
- Write DataFrame or SQL computation!
Integrates well with batch functionality!
- Output to Parquet (for downstream usage)

Structured Streaming Processing



1. Append new row
2. Process new row,
 - update internal state
 - update result
- (3.) Evict oldest rows

Structured Streaming In/Out

Input from

- Apache Kafka
- Files
 - HDFS
 - S3
- Socket

Output to

- Apache Kafka
- File
- Foreach sink
- Console sink (for testing)
- Memory sink (for debugging)

- Developed at TU Berlin, Apache project since 2014
- Big data processing engine: distributed and scalable streaming dataflow engine
- Exploits data streaming and in-memory processing and iteration operators to improve performance
- Seen as 4th Generation IFP system because:
 - Iterative algorithms
 - Internal optimization mechanisms
 - Uses Lambda architecture model
 - Hybrid programming architecture: allows simultaneous batch and real-time runs.



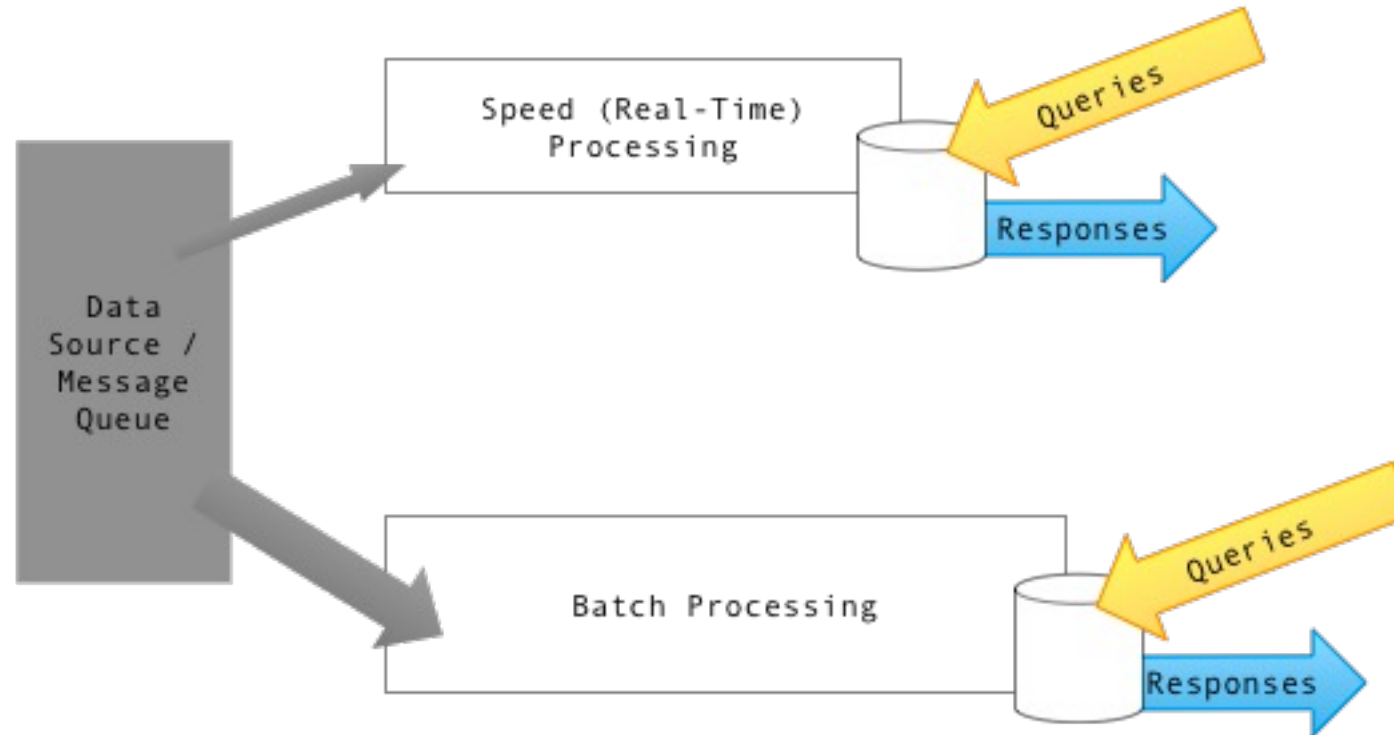
Who uses Apache Flink?

- **Alibaba** uses a Flink-based system to optimize search rankings in real-time
- **Ericsson** used Flink to build a real-time anomaly detector over large infrastructures using machine learning
- **Huawei** Cloud offers a product called "CloudStream", based on Apache Flink
- **Netflix, Uber, Zalando, ...**

- Several APIs in Java/Python/Scala
 - Dataset API – Batch processing
 - DataStream API – Real-time streaming analytics
 - Table API – Relational Queries
- Domain specific libraries
 - FlinkML: Machine Learning Library for Flink
 - Gelly: Graph library for Flink
- Shell for interactive data analysis

- Batch layer: stores ALL the incoming data in an immutable master dataset and pre-computes batch views on historic data.
- Serving layer: indexes views on the master dataset.
- Real-time processing layer: requests data views depending on incoming queries.

Lambda architecture



Stream Processing Frameworks

G. van Dongen and D. Van den Poel, "Evaluation of Stream Processing Frameworks," in *IEEE Transactions on Parallel and Distributed Systems*, vol. 31, no. 8, pp. 1845-1858, 1 Aug. 2020, doi: 10.1109/TPDS.2020.2978480

- Introduction
- Record-at-a-time streaming
- Declarative, functional streaming
- Declarative, relational streaming

What's next – Processing Graph Data

- Network Theory (Briefly)
- Data representation
- Graph Processing Examples
- Distributed Systems for Graph Processing